

# Indexing techniques for databases

Hans Philippi

February 24, 2025

# Why indexing?

- Suppose you are a police officer
- A suspicious car is passing by at high speed
- You want to check the license plate



# Why indexing?

- Around 10.000.000 cars in the Netherlands
- Query: search a car based on license plate
- Assumptions:
  - A tuple (record) takes 400 bytes
  - A hard disk block contains 16 kbyte, so we have 40 tuples on a block
  - A disk IO takes 5 msec
- Maximum search time (complete table scan)
- $10.000.000 / 40 = 250.000$  disk IO
- Search time : 1250 sec = 21 minutes
- Required search time :  $< 1$  sec

# Memory characteristics (figures from 2023)

## Main memory

- Typical size : 4 – 256 GB
- Access time: 100 nsec ( $10^{-7}$  sec)
- Volatile

## Harddisk

(block size: 2 – 32 kbyte)

- Typical size : 2 – 14 TB
- Access time: 5–10 msec ( $10^{-2}$  sec)
- Some speedup possible with clustering
- Non-volatile

## SSD

(expensive)

- Typical size : 256 GB – 8 TB
- Access time: 0,1 msec ( $10^{-4}$  sec)
- Non-volatile


- Indexing enables a quick table search, based on the value of a specific attribute
- Indexing also supports query processing and optimization
- Indexing supports primary key maintenance and uniqueness constraints (other candidate keys)
- Syntax for SQL DDL:

```
CREATE INDEX Person_dob_ndx  
ON Person (date_of_birth);  
CREATE UNIQUE INDEX Person_ppn_ndx  
ON Person (passport_number);
```

# Indexing techniques


- Two fundamental techniques
  - Indexing based on search trees
  - Indexing based on hashing
- Both techniques are applicable to main memory as well as external memory
- Both techniques deal with block sized memory traffic

# Trees according to computer scientists



# Search trees



- The most well known search tree is the binary search tree
- Maximum number of entries when level =  $n$ ?
- Expected number of pointer chases with  $N$  entries?





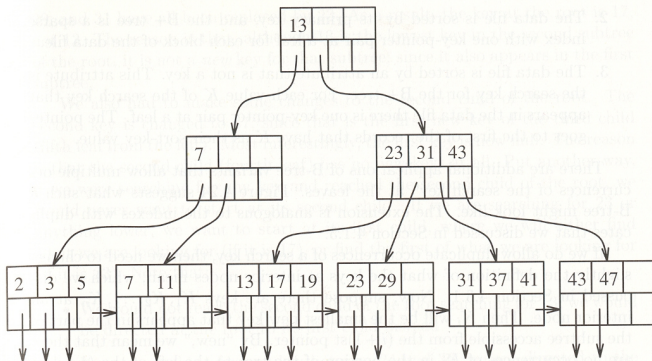
# Search trees: issues

- The most well known search tree is the binary search tree
- Search time is  $O(\log(n))$  for  $n$  entries, when balanced
- Problem: maintaining balance under updates



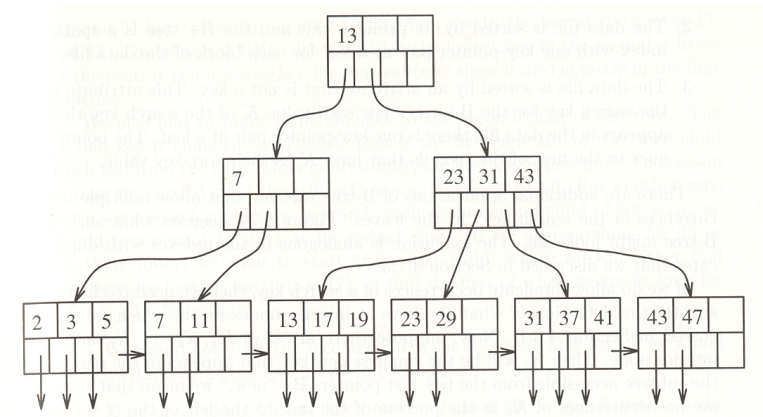
# B-tree

- Standard multiway search tree applied in relational databases
- Sophisticated updating techniques to keep it balanced
- Guarantees at least 50% filling of nodes
- Nodes correspond to disk blocks



# B-tree

- Lowest level contains all attribute values and pointers to corresponding tuples
- Lowest level contains sibling pointers supporting range queries

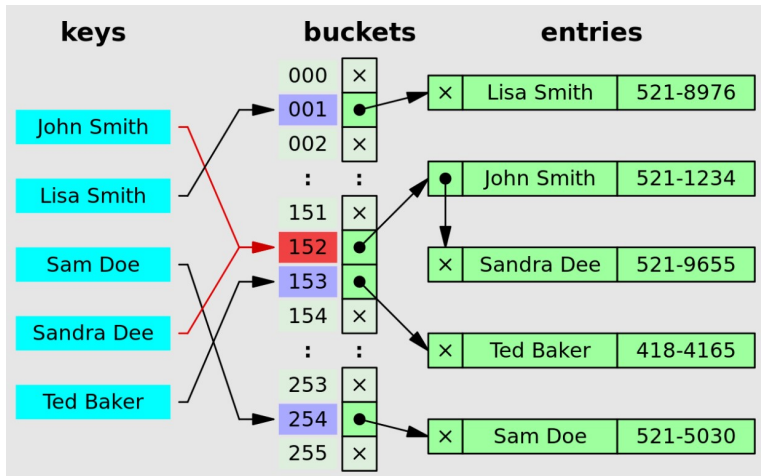


## B-tree: an example

- Attribute value: 4 byte integer
- Pointer: 8 bytes
- Block size: 16 kbyte
- Content: 683 – 1365 entries per block
- 2 levels: minimum nr of entries = 466000
- 3 levels : minimum nr of entries = 318 million
- 4 levels : minimum nr of entries = 217 billion
- Number of pointer traces is limited by  $\lceil k \log(n) \rceil$ , with  $k = 683$
- Search time in our example:  $\ll 1$  sec

- Memory reservation of  $N$  buckets: virtual addresses  $0..N-1$
- Hashfunction  $f$ 
  - Domain: all possible attribute values
  - Codomain:  $0..N-1$
- The hash function calculates the bucket address from the current attribute value
- Hashfunction  $f$  should distribute the addresses evenly
- More info: [https://en.wikipedia.org/wiki/Hash\\_function](https://en.wikipedia.org/wiki/Hash_function)

# Hash table



Source: [https://en.wikipedia.org/wiki/Hash\\_function](https://en.wikipedia.org/wiki/Hash_function)

# Final words

- Hash indexing has a theoretical advantage: one disk access versus  $k \log(n)$  for B-tree
- Hash indexing has a fundamental disadvantage: range queries are not supported ...
- ... while B-trees support range queries by horizontal links on the lowest level
- The  $k$  of  $k \log(n)$  is very large, so  $k \log(n)$  hardly exceeds 3 ...
- ... while the root of the B-tree is (and possibly the second level nodes are) often kept in main memory
- Overall, the B-tree is the winner