

# Indexing in databases


# The problem



# The problem

- ▶ Around 10.000.000 cars in the Netherlands
- ▶ Query: search a car based on license plate
- ▶ Assumptions:
  - A tuple (record) takes 400 bytes
  - A hard disk block contains 16 kbyte, so we have 40 records on a block
  - A disk IO takes ~5 msec
- ▶ Maximum search time (complete table scan)
  - $10.000.000 / 40 = 250.000$  disk IO
  - 1250 sec  $\approx$  21 minutes
- ▶ Required search time :  $< 1$  sec

# Hard disk



# Main memory vs harddisk

## ▶ Main memory

- Typical size : 4 – 256 GB
- Access time:  $\sim 100$  nsec ( $10^{-7}$  sec)
- Volatile

## ▶ Harddisk

- Typical size : 1 – 14 TB
- Access time: 5–10 msec (without clustering)
- Non-volatile

## ▶ SSD

- Typical size : 128 GB – 4 TB (expensive)
- Access time:  $\sim 0,1$  msec ( $10^{-4}$  sec)
- Non-volatile

## ▶ Unit of traffic: block (2 – 32 kbyte)

# The solution: indexing

- ▶ Indexing enables a quick table search, based on the value of a specific attribute
- ▶ Indexing also supports query processing and optimization
- ▶ Indexing (automatically) supports primary key maintenance and uniqueness constraints (other candidate keys)
- ▶ Syntax for SQL DDL:

```
CREATE INDEX Person_dob_ndx
ON Person (date_of_birth);
CREATE UNIQUE INDEX Person_ppn_ndx
ON Person (passport_number);
```




# Indexing: how does it work?

Two fundamental techniques

1. Search tree
2. Hash index


Both techniques applicable to main memory as well as external memory

# Tree according to computer scientists






# Binary search tree




# Intermezzo

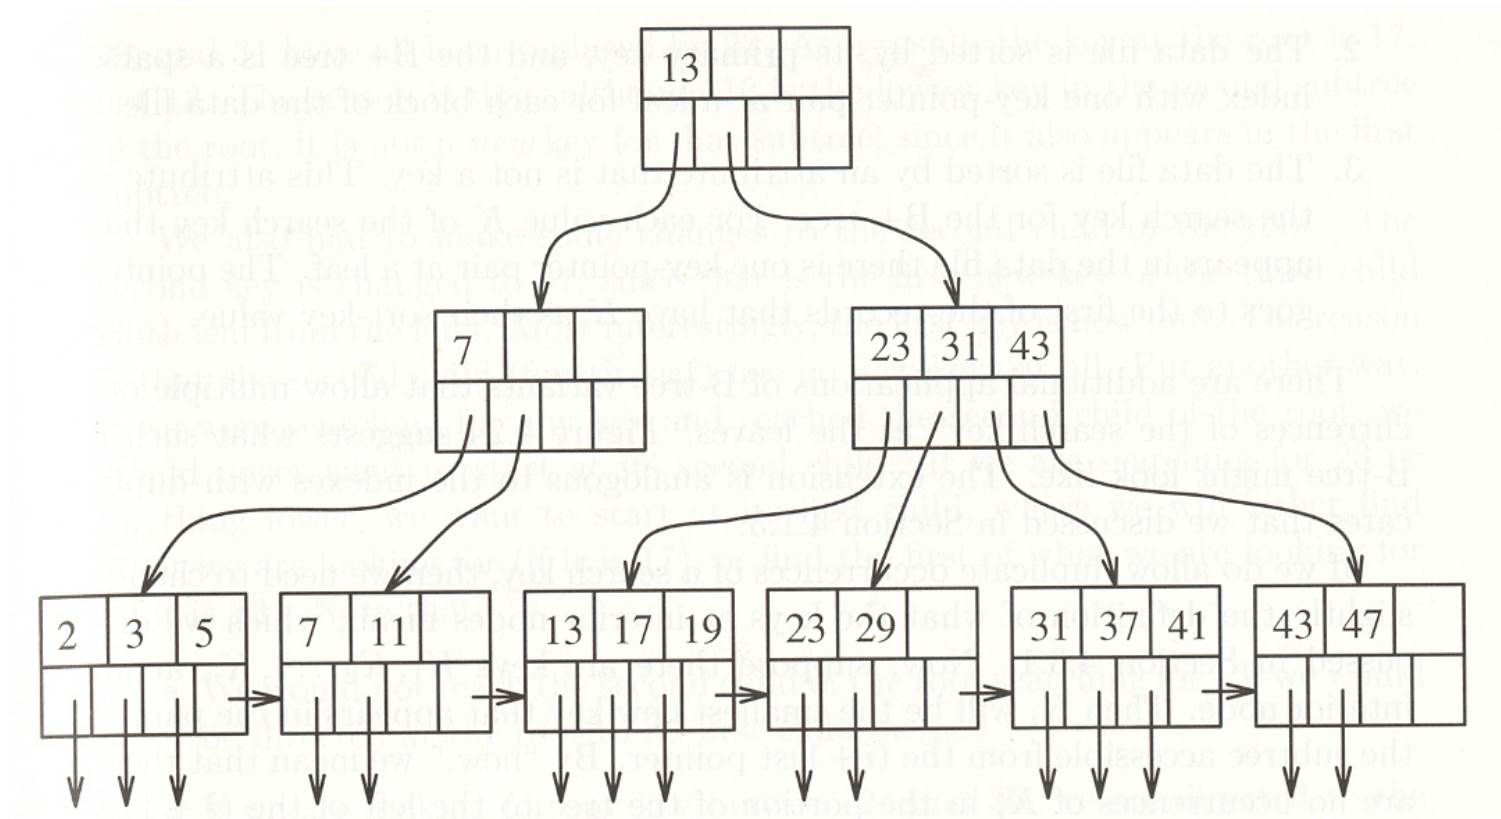


- Add: 2 18 9 5
- How many nodes could a tree with 3 levels contain?
- How many nodes could a tree with N levels contain?
- How many steps does it (roughly) take to fulfill a search action in a tree with M nodes?
- Add: 20 30 40 60 50

# Binary search tree: balancing



# B-tree (main memory & harddisk)



Update algorithm guarantees 50% filling of nodes

# B-tree: characteristics

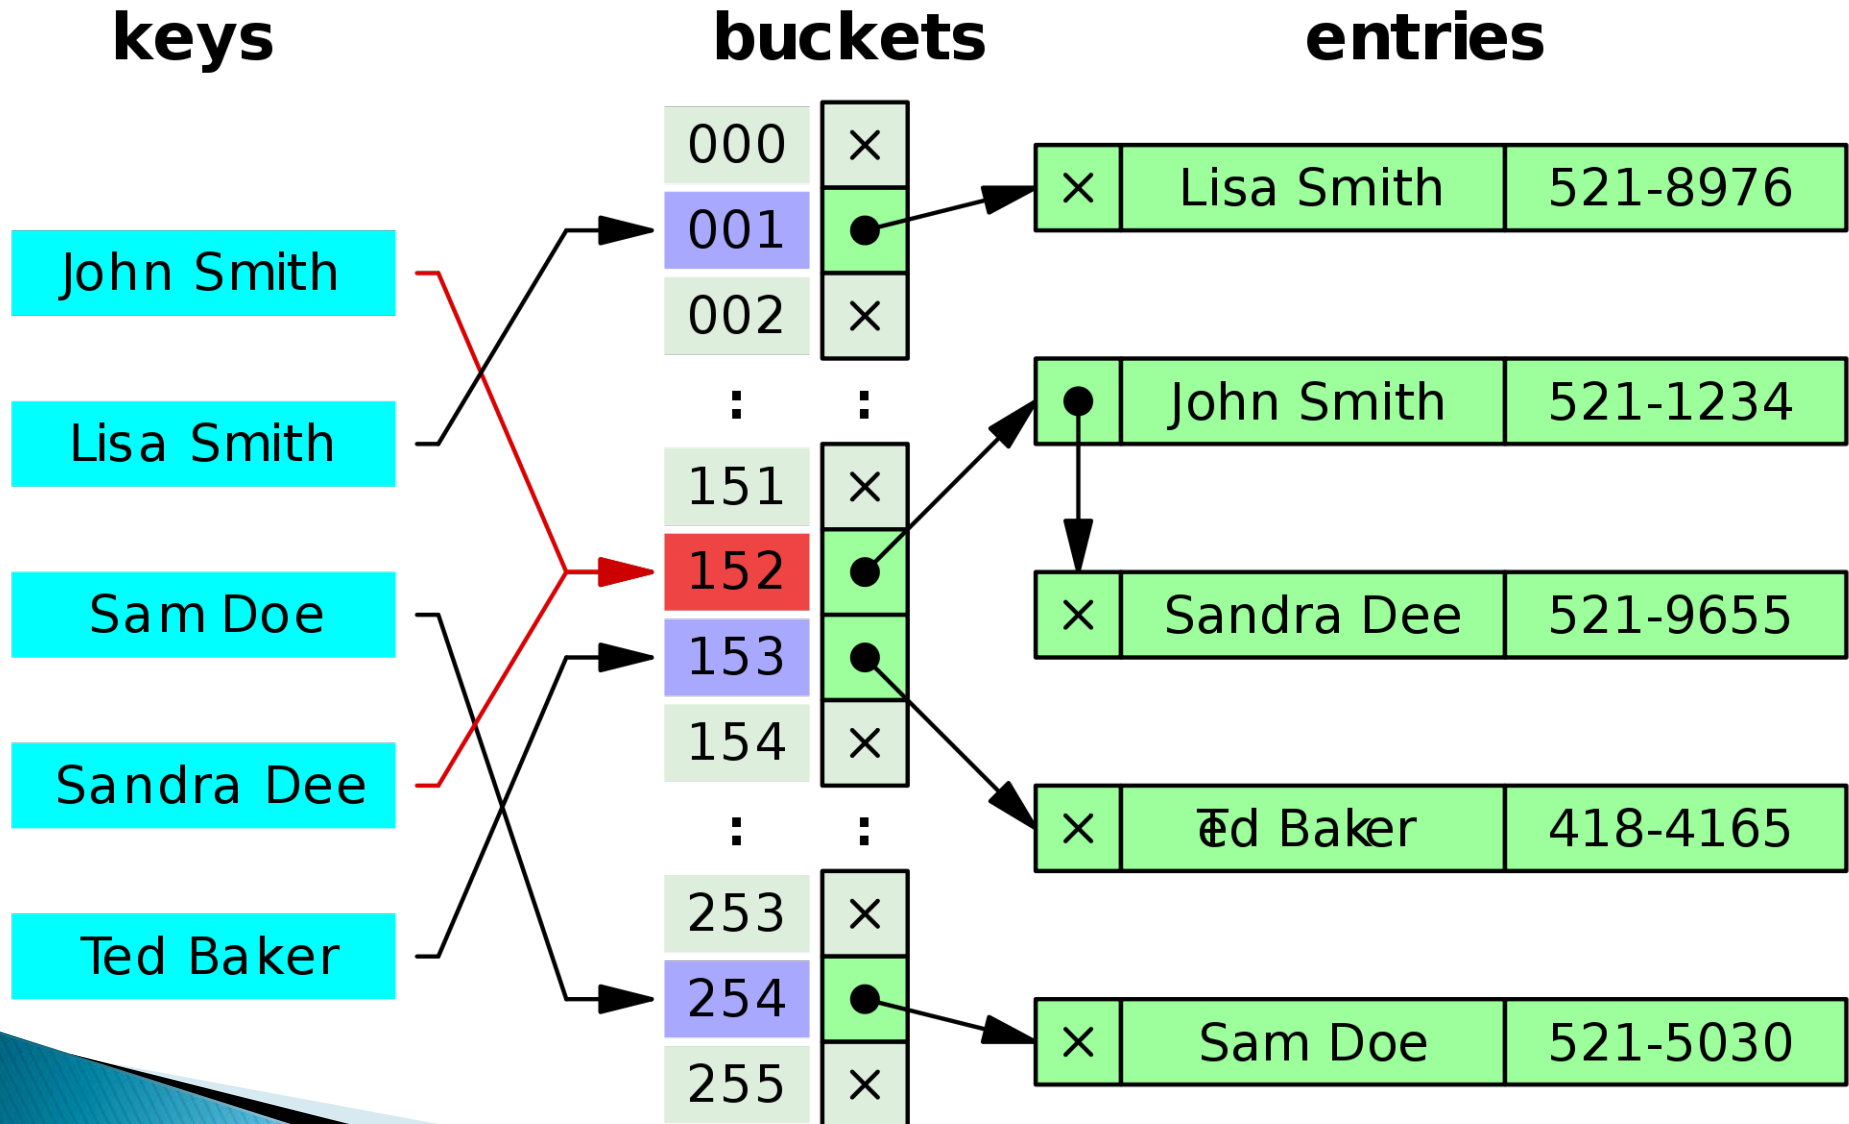
- ▶ Field: 4 byte integer
- ▶ Pointer: 8 bytes
- ▶ Block size: 16 kbyte
- ▶ Content: 683 – 1365 entries per block
- ▶ 2 levels: minimum nr of entries = 466489
- ▶ 3 levels : minimum nr of entries = 318 million
- ▶ 4 levels : minimum nr of entries = 217 billion



# Hash table

- ▶ Memory reservation of  $N$  buckets: addresses  $0..N-1$
- ▶ Hashfunction  $f$ 
  - Domain: all possible attribute values
  - Codomain:  $0..N-1$
- ▶ Hopefully,  $f$  distributes the values neatly

# Hash table



# Indexing: final words

- ▶ Hash indexing has a theoretical advantage: one disk access versus  $k \log(M)$  for B-tree
- ▶ Hash indexing has a fundamental disadvantage: range queries are not supported
- ▶ The  $k$  of  $k \log(M)$  is very large, so  $k \log(M)$  hardly exceeds 3 ...
- ▶ ... while the root of the B-tree (and possibly the second level nodes) are often kept in main memory
- ▶ Overall, the B-tree is the winner