


# Transaction Processing Concurrency

Hans Philippi

January 9, 2020

# Transactions

Transactions are about a group of database operations (inserts, deletes, updates) that should be regarded as one logical unit and one atomic action:



Applications often require the possibility to execute transactions in a concurrent mode:

The Seattle-based company said customers ordered 34.4 million items during [Prime Day](#), or to put it another way, 398 items per second. In the U.S. alone it sold 47,000 TVs, marking a 1300 percent increase on last year; 41,000 Bose headphones, compared to eight the previous Wednesday; and 14,000 iRobot Roomba 595 Pet Vacuum Cleaning Robots, compared to one the

# Transaction concept

*Definition:*

## Atomic database action

$R(x)$ : read data element  $x$

$W(x)$ : write data element  $x$

data element: attribute, tuple, block, table

*Definition:*

A transaction is a collection of atomic database actions considered to be one logical unit, with respect to:

- concurrency
- recovery

# Transaction concept

## Termination of a transaction

- *commit*: positive, complete, irreversible, persistent; effects to real world will become visible
- *abort* or *rollback*: negative, nihil; no effects to real world; no partial results visible

## ACID-properties:

**A** tomicity

**C** onsistency preservation

**I** solation

**D** urability

# Schedules

The (concurrent) execution order of a sequence of database operations generated by a set of transactions is represented by a *schedule* or *history*

*Example:*

an interleaved schedule  $S1$  and a serial schedule  $S2$

$S1$		
T1	T2	T3
	W(y)	
		R(x)
W(x)		
	W(z)	
		R(y)

↓ time

$S2$		
T1	T2	T3
	W(y)	
	W(z)	
		R(x)
		R(y)
W(x)		

*Problem 1: lost update*

$S3$

T1	T2
R(x)	
	R(x)
W(x)	
	W(x)

# Inconsistent retrieval

*Problem 2: inconsistent retrieval*  
*S4*

T1	T2
	$R(x_1)$
	$R(x_2)$
	$\vdots$
	$R(x_m)$
$R(x_m)$	
$W(x_m)$	
$R(x_{m+1})$	
$W(x_{m+1})$	
	$R(x_{m+1})$
	$\vdots$
	$R(x_n)$



## Formalizing correctness

- serial schedules (schedules without concurrency) are correct
- two schedules are equivalent iff their effects visible to the outside world are exactly the same
- an interleaved schedule is correct iff it is equivalent to a serial schedule; such a schedule is called serializable

- *Intuition:* two atomic database actions are conflicting if the effects of reversing the relative order of their execution are visible to the outside world
- *Definition:* two atomic database actions are conflicting iff they affect the same data element and at least one of the operations is a write
- *Definition:* two schedules are equivalent iff they order all the conflicts the same way

# Correctness criterion

S5			S6		
T1	T2	T3	T1	T2	T3
		W(x)			W(x)
R(x)					R(z)
		R(z)	R(x)		
	W(x)		W(y)		
W(y)				W(x)	
	W(y)			W(y)	

# Correctness criterion

S5			S6		
T1	T2	T3	T1	T2	T3
		W(x)			W(x)
R(x)					R(z)
		R(z)	R(x)		
	W(x)		W(y)		
W(y)				W(x)	
	W(y)			W(y)	

S5 orders all the conflicts the same way S6 does

... so S5 is equivalent to S6 ...

... and S6 is serial ...

... so S5 is correct: *serializable*

<i>S7</i>				<i>S8</i>			
T1	T2	T3	T4	T1	T2	T3	T4
		W(x)				W(x)	
	R(x)				R(x)		
W(z)							R(z)
			R(z)				W(y)
			W(y)	W(z)			
		W(y)				W(y)	
	R(y)				R(y)		
W(y)				W(y)			
	W(x)				W(x)		

is *S7* serializable ?

is *S8* serializable ?

# Testing serializability

So, testing serializability means  
comparing with ... serial schedules of  $n$  transactions?

*Definition:* a precedence graph  $G(S)$  is a directed graph connected to a schedule  $S$  such that:

- the set of nodes is the set of transactions in  $S$
- there exists an edge  $T_i \rightarrow T_j$  in  $G$  iff  
there is a conflicting pair of operations  $o_i, o_j$  in  $S$   
such that  $o_i$  occurs before  $o_j$

*Theorem:*

a schedule  $S$  is serializable  $\Leftrightarrow G(S)$  is acyclic



# Enforcing serializability: 2PL

## *Two phase locking (2PL)*


Scheduling technique based on locking principle

- before executing an atomic operation, the requesting transaction has to acquire an exclusive lock on the data element; if the data element is locked by another transaction, the requesting transaction has to wait until release of the lock by the other transaction
- after executing the operation, the scheduler has to release the lock


# Enforcing serializability: 2PL

- the transaction is not allowed to acquire new locks after the release of a lock; so locking is executed in two consecutive phases: the growing phase and the shrinking phase (with respect to the number of locks held by the transaction)

NUMBER  
OF LOCKS



2PL-BEHAVIOUR



NON-2PL-BEHAVIOUR

# Enforcing serializability: 2PL

*Theorem:*

$S$  is a 2PL-schedule  $\Rightarrow S$  is serializable

# Enforcing serializability: 2PL

## Refinement

- concurrent execution of read-operations should be allowed
- exclusiveness of locks is only required in case of conflicting operations
- we distinguish read-locks and write-locks (or exclusive locks)

## Lock compatibility matrix

	read	write
read	yes	no
write	no	no

Note that 2PL behaviour is enforced on *every* transaction separately!

# Enforcing serializability: 2PL

*Problem:* deadlock

*Solution:*

- prevention
- detection by time out
- detection by wait-for-graphs

# Enforcing serializability: timestamping

## Alternative to 2PL

- Every transaction receives a timestamp at birth
- Timestamps are emitted in some order (typically a counter or system time)
- Conflicting operations are executed in timestamp order
- Application: ad hoc concurrency control in loosely coupled databases

# Enforcing serializability: timestamp ordering principle

Timestamp rule:

For each conflicting pair of operations  $o_i, o_j$  in  $S$ ,  
if  $o_i < o_j$  then  $ts(T_i) < ts(T_j)$

where  $o_i < o_j$  denotes time order

Implementation:

timestamp manager maintains a table of recently used data entries

# Enforcing serializability: timestamping

*Theorem:*

If schedule  $S$  obeys the timestamp rule  
then  $S$  is serializable



# Enforcing serializability: other methods

Optimistic concurrency control by a posteriori validation

no further details