

Index construction and MapReduce

Hans Philippi
(partially based on the slides from the Stanford course on IR)

March 30, 2023

MapReduce: an example of NoSQL data management

- Data model: key-value pairs
- Massive parallelism on ...
- ... large amounts of commodity hardware
- Case: index building for text collections, especially for the Web

Index for text collections

In general for document collections:

INPUT: a term

OUTPUT: the corresponding postings list, i.e. all documents in the text collection containing this term

Specifically for the Web:

INPUT: a term

OUTPUT: the corresponding postings list, i.e. all URLs of web pages containing this term

Creating postings lists

Input : document collection $\langle \text{docid}, \text{text} \rangle$

$\langle 2013, \text{"de dag die je wist dat zou komen is eindelijk hier"} \rangle$

$\langle 1980, \text{"de do do do, de da da da"} \rangle$

$\langle 1971, \text{"jaren komen en jaren gaan"} \rangle$

$\langle 1994, \text{"we komen en we gaan"} \rangle$

Output : a set of *postings lists* for this collection of documents

...

$\langle \text{"dag"}, [2013] \rangle$

$\langle \text{"de"}, [1980, 2013] \rangle$

$\langle \text{"die"}, [2013] \rangle$

$\langle \text{"do"}, [1980] \rangle$

$\langle \text{"en"}, [1971, 1994] \rangle$

$\langle \text{"gaan"}, [1971, 1994] \rangle$

...

$\langle \text{"komen"}, [1971, 1994, 2013] \rangle$

...

INPUT: a term

OUTPUT: the corresponding postings list, i.e. all occurrences of this term in the text collection containing this term.

Two steps:

- Use a tree structure (B-tree, suffix tree) that connects a term to the corresponding postings list
- Return the postings list

Postings list sorted

Query = $term_1$ AND $term_2$

- 1 locate postings list p_1 for $term_1$
- 2 locate postings list p_2 for $term_2$
- 3 calculate the intersection of p_1 and p_2 by list merging

$term_1 \Rightarrow$	1	3	7	11	37	44	58	112	...
$term_2 \Rightarrow$	2	4	11	25	44	54	55	58	...

Index construction: two approaches

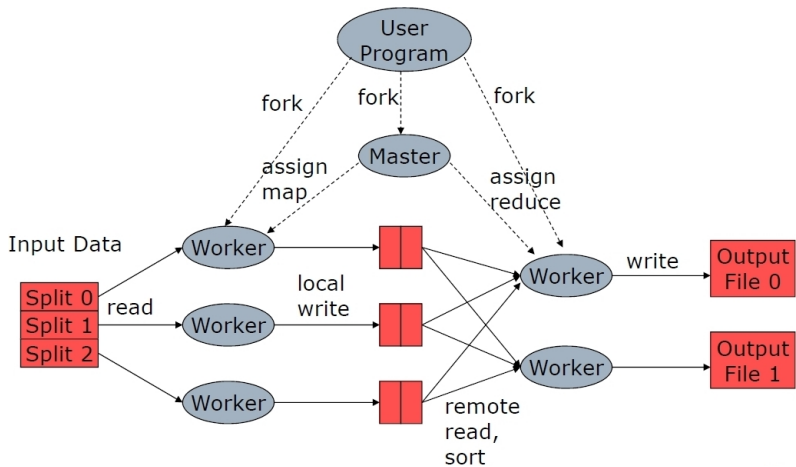
- Algorithms dealing with limited main memory, based on external sorting. Output of sorting phase enables index building.
- Index building based on MapReduce: generic architecture for and approach to large scale parallelism

- Framework for massively parallel computing
- Roots in Google environment (indexing, PageRank)
- Based on commodity hardware
- Two sets of machines involved in parallel processing: *Map* workers and *Reduce* workers
- Robustness by replication of data in file system
- Generic, based on *Map* and *Reduce (Fold)* from functional programming
- Several implementations, Hadoop is the most well known

MapReduce: the Map

- Basic data structure is key-value pair $\langle k, v \rangle$
- Input is split into disjoint chunks, containing collections of key value pairs
- Each Map worker works autonomous from other map workers (“shared nothing”)
- Each Map worker scans it's own input chunk once
- Each Map worker does one uniform calculation on each key-value pair
- The output of each Map worker is a set of key-value pairs: zero, one or more
- The structure of the resulting key-value pairs is generally different from the input pairs
- The output results of all Map workers are collected for further processing in the Reduce phase

MapReduce computing



MapReduce: the Reduce

- The output results of all Map workers are grouped on the key values and assigned to the reduce workers
- All related key value pairs will be processed by one Reduce worker
- Each Reduce worker works autonomous from other Reduce workers (shared nothing)
- The output results of all Reduce workers together are the result of the calculation

MapReduce example

Example: *word count*

Input: a collection of documents

Output: the words in the documents with their frequency

- Map $\langle docid, text \rangle$:
for each word w in $text$
 $emit(\langle w, 1 \rangle)$;
- Reduce $\langle w, vlist \rangle$:
 int $sum = 0$;
 for each v in $vlist$
 $sum ++$;
 $emit(\langle w, sum \rangle)$;

MapReduce example

Input to Map-workers:

< 2013, "de dag die je wist dat zou komen is eindelijk hier" >
< 1980, "de do do do, de da da da" >
< 1971, "jaren komen en jaren gaan" >
< 1994, "we komen en we gaan" >

Output from Map workers:

<"de", 1 >
<"dag", 1 >
<"die", 1 >
...
<"gaan", 1 >

MapReduce example

... then comes the invisible step ...

... which could be characterized as a "GROUP BY key" ...

MapReduce example

Input to Reduce-workers:

$\langle \text{"de"}, [1] \rangle$

...

$\langle \text{"komen"}, [1, 1, 1] \rangle$

...

$\langle \text{"gaan"}, [1, 1] \rangle$

...

Output:

$\langle \text{"de"}, 1 \rangle$

...

$\langle \text{"komen"}, 3 \rangle$

...

$\langle \text{"gaan"}, 2 \rangle$

...

MapReduce example

Observations:

- The input pairs will be processed by different Map-workers
- Behind the scenes (invisible step), all emitted pairs with the same key are grouped together (after the Map phase and before the Reduce phase)
- The *grouping phase* includes concatenation of all the values corresponding to the same key
- In our example: in the grouping phase: three times $\langle \text{"komen"}, 1 \rangle$ becomes $\langle \text{"komen"}, [1, 1, 1] \rangle$

INTERMEZZO MapReduce example

Do you have any suggestions for optimization of the MapReduce program from the example on slide 12?

MapReduce computing: early combining

- Word count could be optimized by doing some aggregation in the Map phase
- Instead of k repetitions of $emit(< w, 1 >)$; do $emit(< w, k >)$;
- Adapt the Reduce program (how?)
- In general, this idea is applicable if the reduce function is commutative and associative (e.g. sum, max)
- Early combining often requires a setup of local datastructures and a final emit
- Our convention: for writing pseudo code, use functions *Init_Map()* and *Finalize_Map()*

Word count speed up

- `Init_Map()`:
Create a dictionary `D (word, freq)`;
- `Map < docid, text >`:
for each word `w` in `text`
add `w` to `D`;
- `Finalize_Map()`:
for each entry `(word, freq)` in `D`
`emit(< word, freq >)`;
- `Reduce < w, vlist >`:
int `sum = 0`;
for each `v` in `vlist`
`sum += v`;
`emit(< w, sum >)`;

Let's do it

Schrijf pseudocode voor Map en Reduce voor een collectie tupels van de vorm $\langle g, v \rangle$ die de volgende SQL query representeert:

```
SELECT g, SUM(v) FROM Input
WHERE v >= 100
GROUP BY g
HAVING SUM(v) >= 10000
```

Let's do it again

Schrijf pseudocode voor Map en Reduce voor een collectie tupels van de vorm $\langle g, v \rangle$ die de volgende SQL query representeert:

```
SELECT g, SUM(v) FROM Input
WHERE v >= 100
GROUP BY g
HAVING SUM(v) <= 10000
```

MapReduce: references

- <http://infolab.stanford.edu/~ullman/mmds/ch2.pdf>