

NoSQL and the idea of scaling up and out Databases



What lies beyond Relational

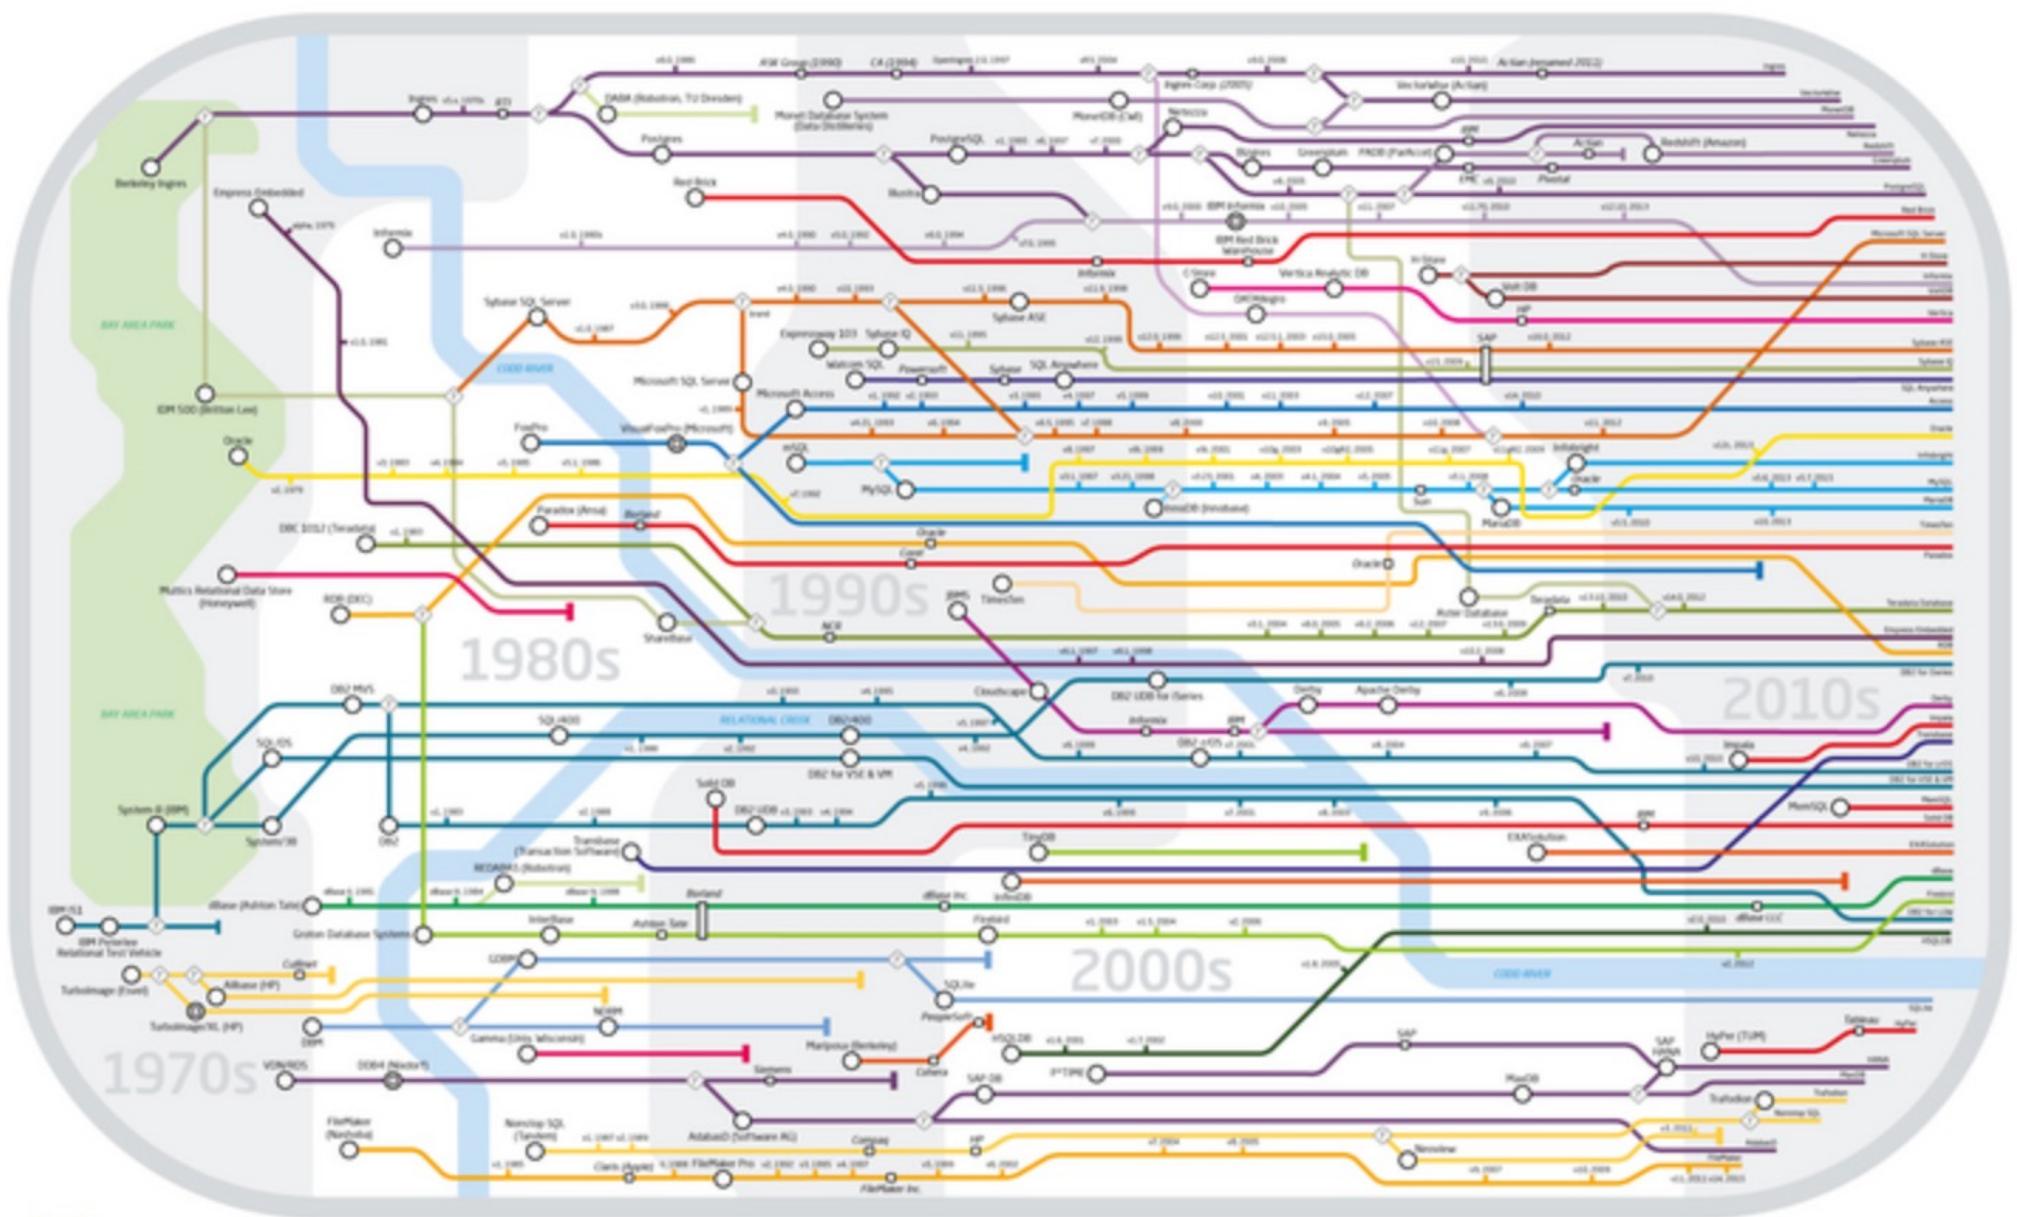
Prof. Yannis Velegrakis
Utrecht University

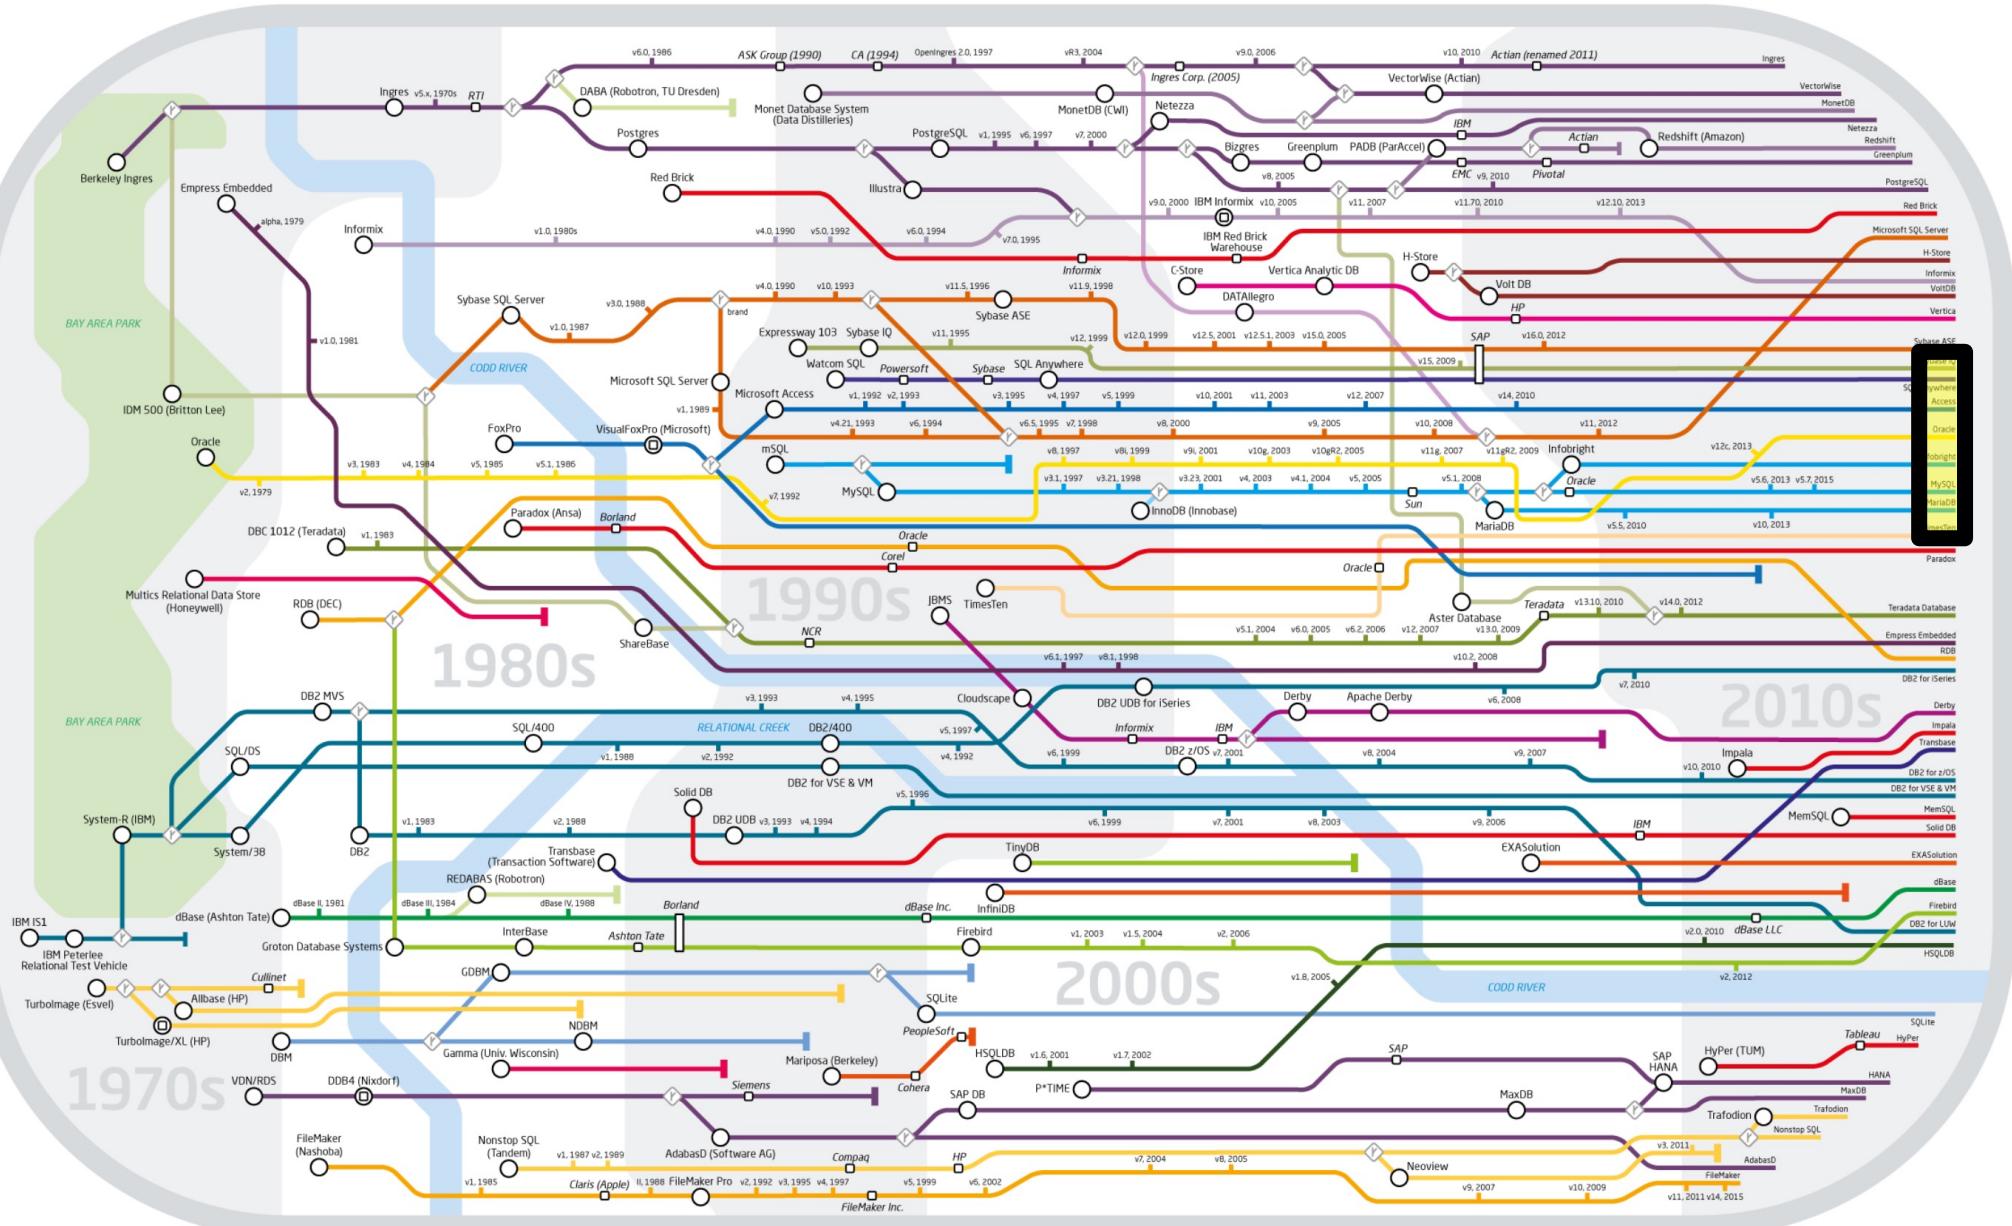
i.velegrakis@uu.nl
<https://velgias.github.io>

Disclaimer:

Slides courtesy of

- Martin Fowler
- Michalis Petropoulos
- Bill Howie
- Yannis Velegrakis









The Need for Data Storage

How do we store data?





Why would I want a database?

What problem do they solve?

1. **Sharing**
Support concurrent access by multiple readers and writers
2. **Data Model Enforcement**
Make sure all applications see clean, organized data
3. **Scale**
Work with datasets too large to fit in memory
4. **Flexibility**
Use the data in new, unanticipated ways

How do we store data?

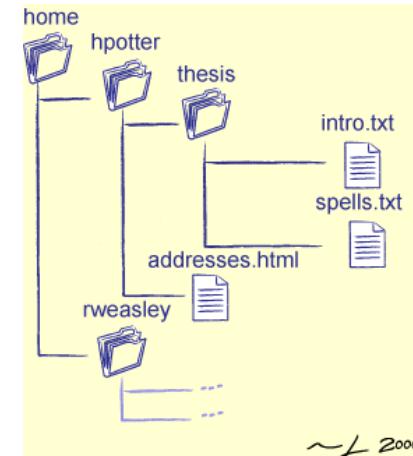
Line P Cruise - GeoMICS (May 2012)

2 Nutrients analyzed on board Thompson by members of the Ingalls and Devol Labs (Laura Truxal, Davey French, Katherine Heal
3 ~Water sampled from CTD Niskin bottles unless otherwise indicated.

Station P8 Nutrients.

	Depth (m)	Conc NO2 (nM)	Conc NH4 (nM)	
8	5	0	35.67	NH4 Analysis - OPA ammonia assay
9	35	125	181.89	Reference:
10	40	110		Holmes, R. M., A. Aminot, R. Kerouel, B. A. Hoo
11	45	165		Kerou & Aminot (1997). Fluorometric determin
12	50	125		NO2 Analysis - NO2 Colorimetric Assay
13	60	290		Reference:
14	70	445	0	Grashoff, Kremling, Erhard (1999). Methods o
15	85	0		
16	105	0		
17	300	0	0	
18	GoFlo 0055	70	455	16.06
19	GoFlo 0052	70	445	3.51
20				
21				
22	Station P6 Nutrients			
23		Conc NO2	Conc NH4	

What is the *data model*?



ANNOTATIONSUMMARY-COMBINEDORFANNOTATION16_Phaeo_genome

##query	length	COG hit #1	e-value #1	identity #1	score #1	hit length #1	description #1
chr_4[480001-580000].287	4500						
chr_4[560001-660000].1	3556						
chr_9[400001-500000].503	4211	COG4547	2.00E-04	19	44.6	620	Cobalamin biosynthesis protein
chr_9[320001-420000].548	2833	COG5406	2.00E-04	38	43.9	1001	Nucleosome binding factor SPM
chr_27[320001-404298].20	3991	COG4547	5.00E-05	18	46.2	620	Cobalamin biosynthesis protein
chr_26[320001-420000].378	3963	COG5099	5.00E-05	17	46.2	777	RNA-binding protein of the Puf
chr_26[400001-441226].196	2949	COG5099	2.00E-04	17	43.9	777	RNA-binding protein of the Puf
chr_24[160001-260000].65	3542						
chr_5[720001-820000].339	3141	COG5099	4.00E-09	20	59.3	777	RNA-binding protein of the Puf
chr_9[160001-260000].243	3002	COG5077	1.00E-25	26	114	1089	Ubiquitin carboxyl-terminal hyd
chr_12[720001-820000].86	2895	COG5032	2.00E-09	30	60.5	2105	Phosphatidylinositol kinase and
chr_12[800001-900000].109	1463	COG5032	1.00E-09	30	60.1	2105	Phosphatidylinositol kinase and
chr_11[1-100000].70	2886						
chr_11[80001-180000].100	1523						

What is a Data Model?

Three components:

1. Structures
2. Constraints
3. Operations

Examples

1. Structures

- rows and columns?
- nodes and edges?
- key-value pairs?
- a sequence of bytes?

2. Constraints

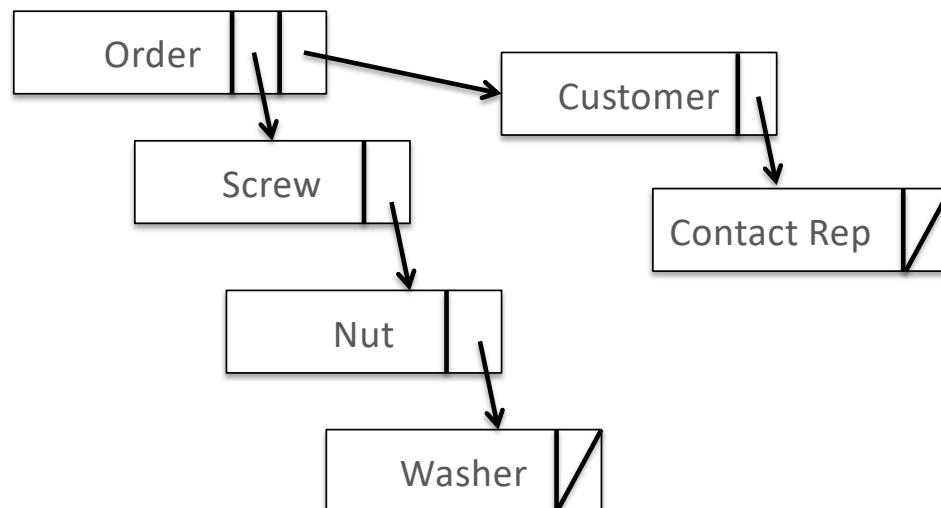
- all rows must have the same number of columns
- all values in one column must have the same type
- a child cannot have two parents

3. Operations

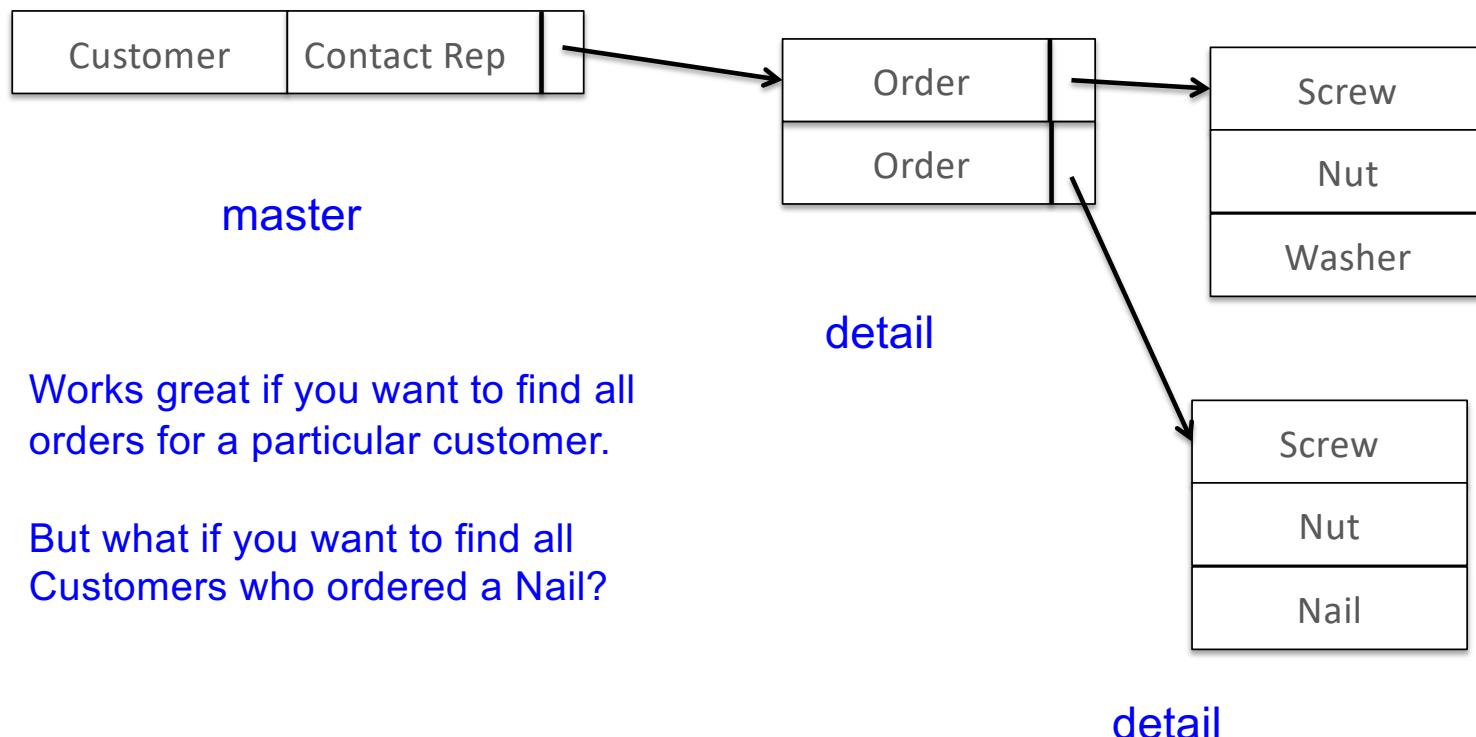
- find the value of key x
- find the rows where column “lastname” is “Jordan”
- get the next N bytes

Database: A collection of information organized to afford efficient retrieval

Historical Example: Network Databases



Historical Example: Hierarchical Databases

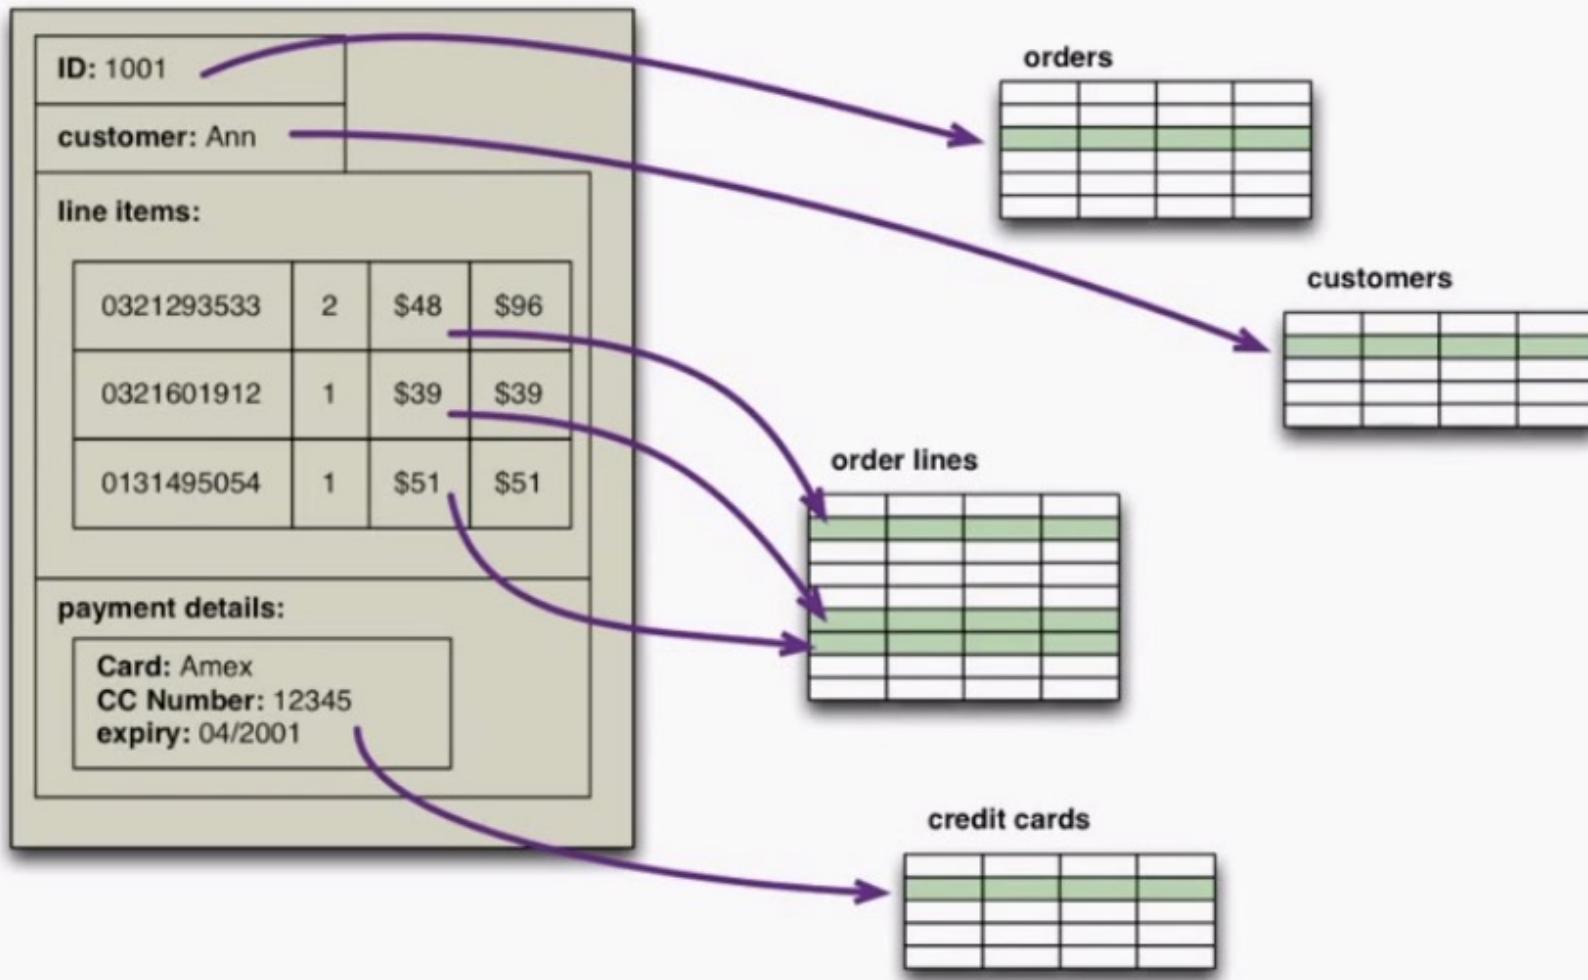


Relational Databases (Codd 1970)

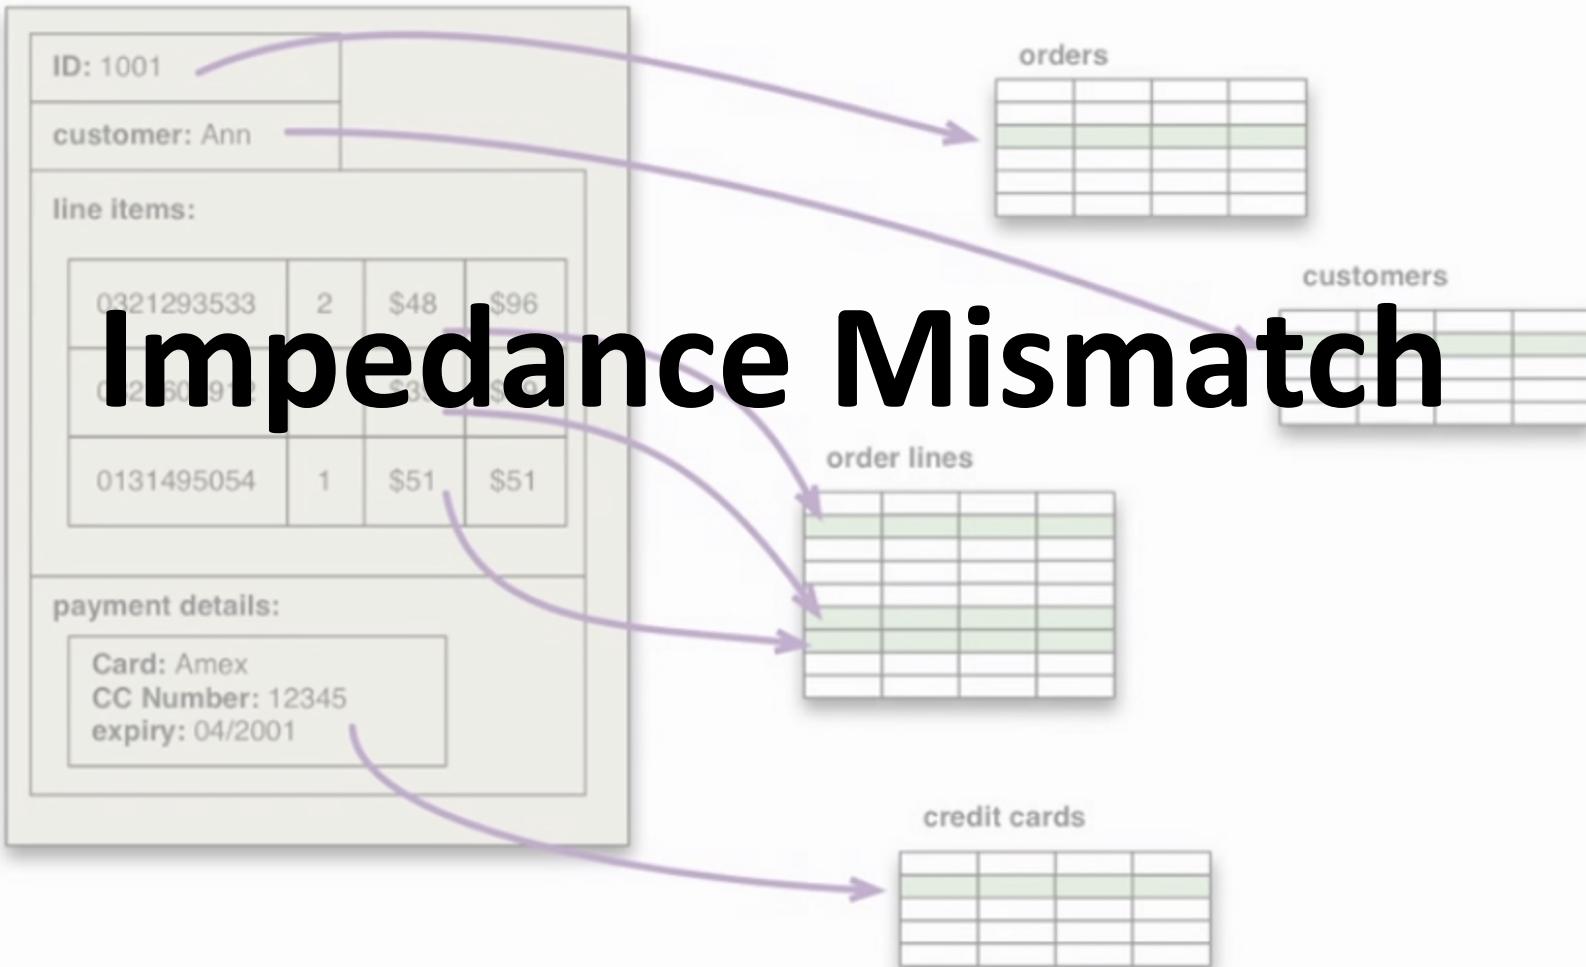
- Everything is a table
- Every row in a table has the same columns
- Relationships are implicit: no pointers

Course	Student Id
CSE 344	223...
CSE 344	244...
CSE 514	255..
CSE 514	244...

Student Id	Student Name
223...	Jane
244...	Joe
255..	Susan



Impedance Mismatch



1980

1990

2000

2010

Rise of the Relational Databases

1980

1990

2000

2010

Rise of the Object Databases

1980

1990

2000

2010

Relational Dominance

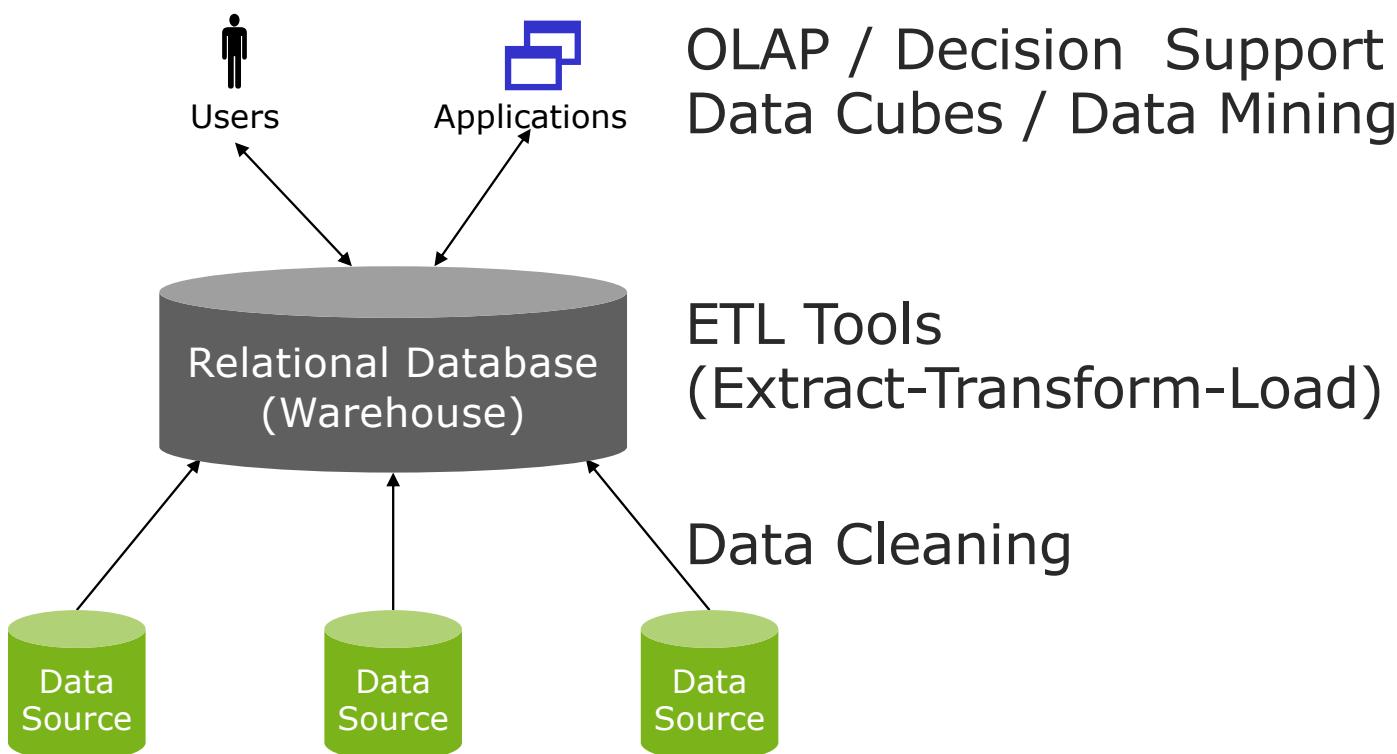


The Need for Integration

Speculative Retailers Web Application

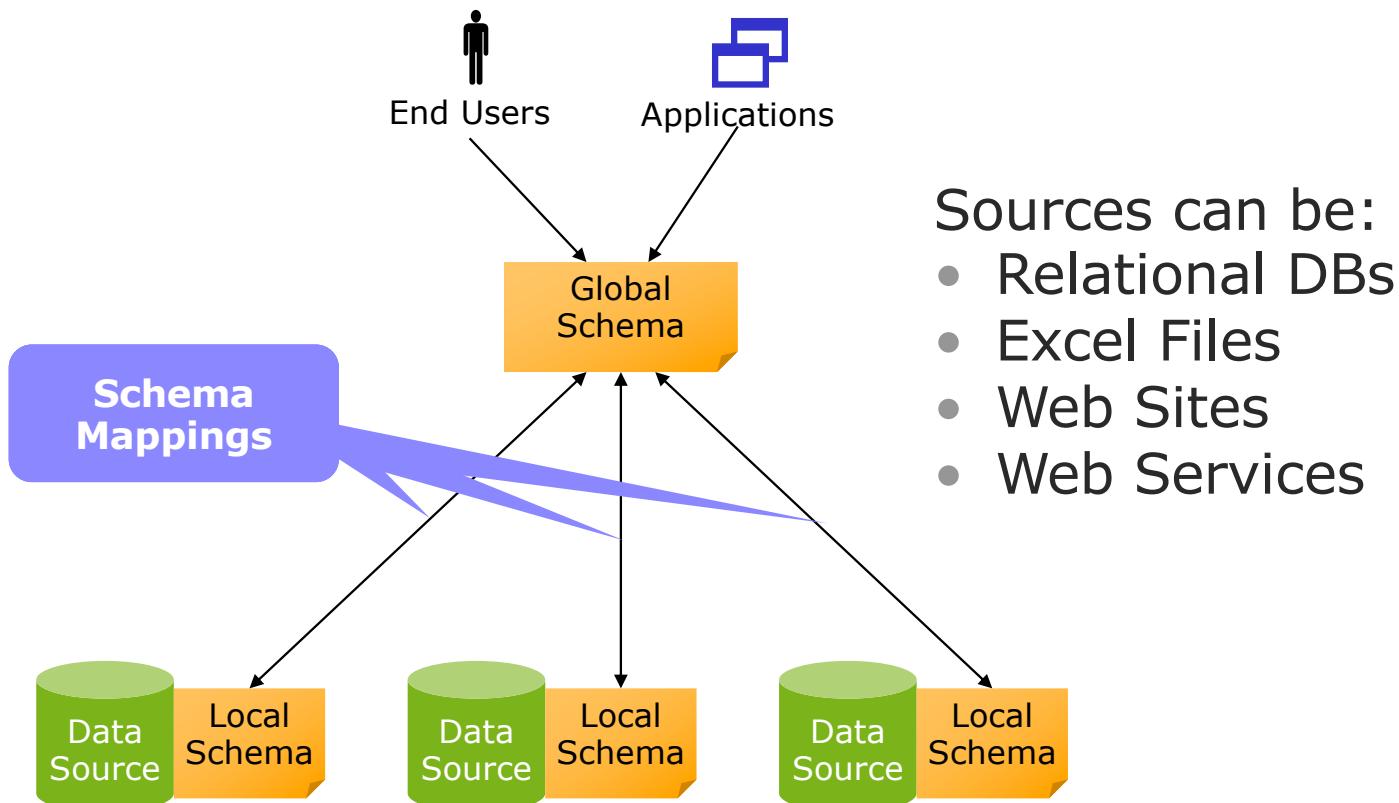


Data Warehouse Architecture

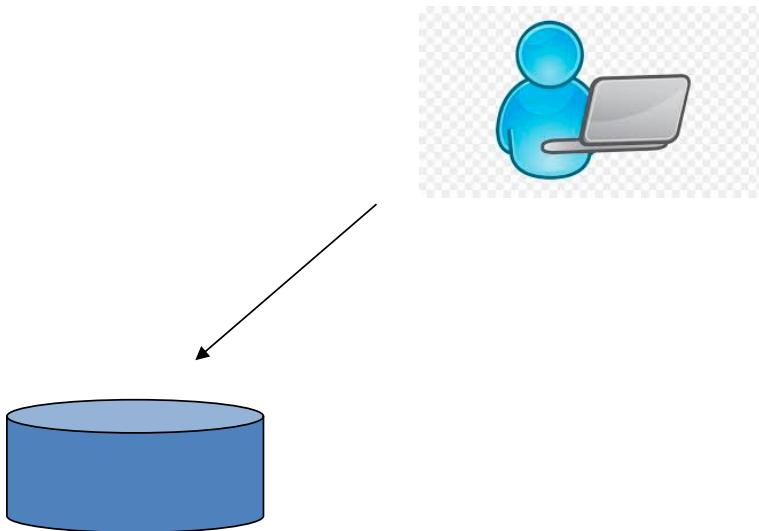


Virtual Integration Architecture

Design-Time

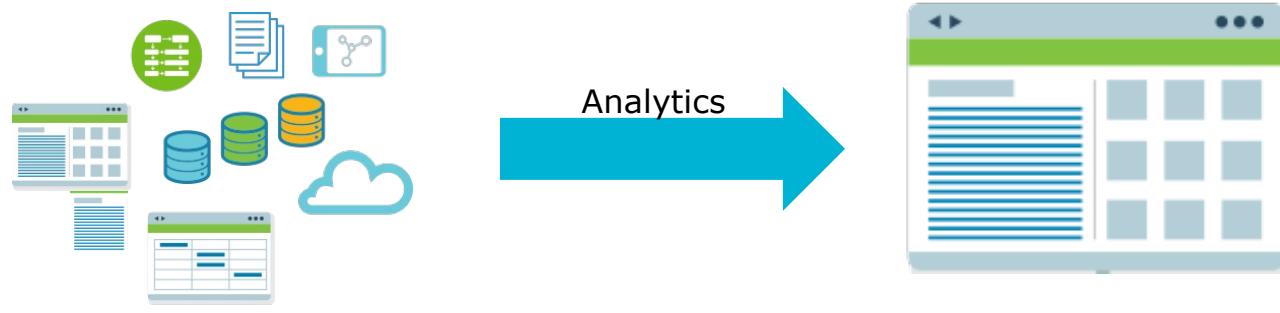


Data Management: The Old Model



Data Analytics: The New Model

Extracting Value from Data





Real Life Data Use Cases

- Education
 - Correlate academic effectiveness with various factors
- Healthcare
 - Preventive & Personalized
- Urban Planning
 - Fusion of high fidelity geographical data
- Intelligent Transportation
 - Analysis and visualization of live and detailed road network data
- Environmental Modeling
 - Sensor Networks



Real Life Data Use Cases

- Energy Saving
 - Unveiling Patterns of Use
- Financial Systemic Risk Analysis
 - Analyze contracts and relationships among entities
- Security
 - Social Networks, Financial Transactions, Transportation
- Computer Security
 - Log Info, Data Transfers and Transactions
- Democracy
 - Citizens understand how the government operates



The Need for Scaling Up

Machine Generated Data

Internet Users

Employees / Data Engineers

More Data

2020 This Is What Happens In An Internet Minute

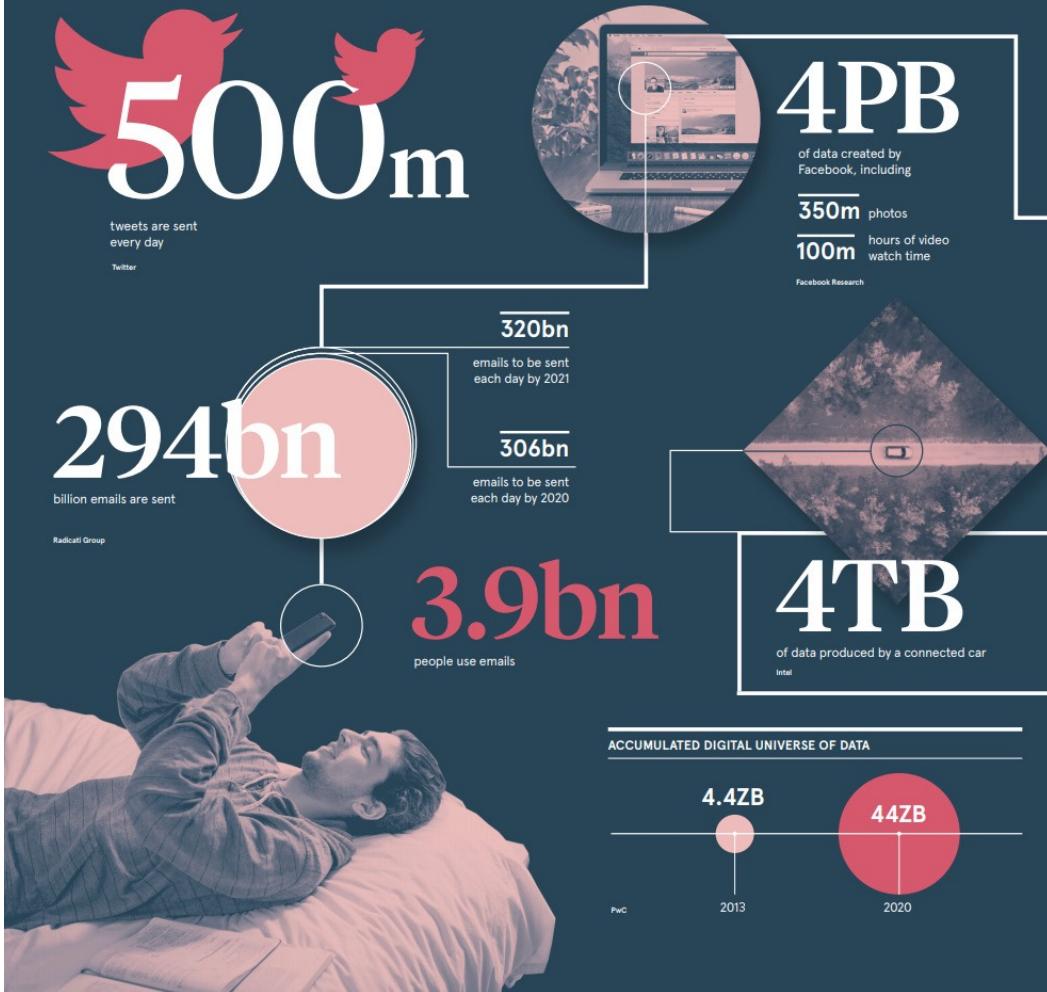


2021 This Is What Happens In An Internet Minute



A DAY IN DATA

The exponential growth of data is undisputed, but the numbers behind this explosion – fuelled by internet of things and the use of connected devices – are hard to comprehend, particularly when looked at in the context of one day



DEMYSTIFYING DATA UNITS

From the more familiar 'bit' or 'megabyte', larger units of measurement are more frequently being used to explain the masses of data

Unit	Value	Size
b bit	0 or 1	1/8 of a byte
B byte	8 bits	1 byte
KB kilobyte	1,000 bytes	1,000 bytes
MB megabyte	1,000 ² bytes	1,000,000 bytes
GB gigabyte	1,000 ³ bytes	1,000,000,000 bytes
TB terabyte	1,000 ⁴ bytes	1,000,000,000,000 bytes
PB petabyte	1,000 ⁵ bytes	1,000,000,000,000,000 bytes
EB exabyte	1,000 ⁶ bytes	1,000,000,000,000,000,000 bytes
ZB zettabyte	1,000 ⁷ bytes	1,000,000,000,000,000,000,000 bytes
YB yottabyte	1,000 ⁸ bytes	1,000,000,000,000,000,000,000,000 bytes

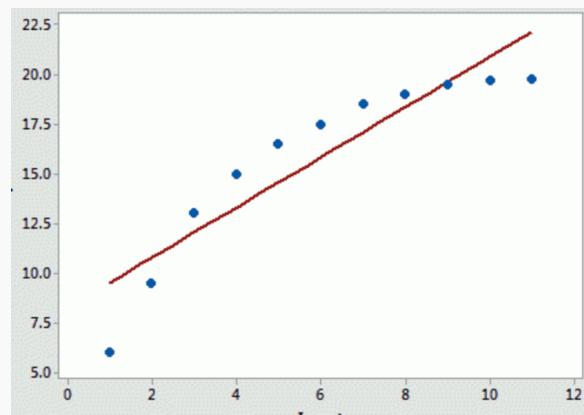
*A lowercase "b" is used as an abbreviation for bits, while an uppercase "B" represents bytes.





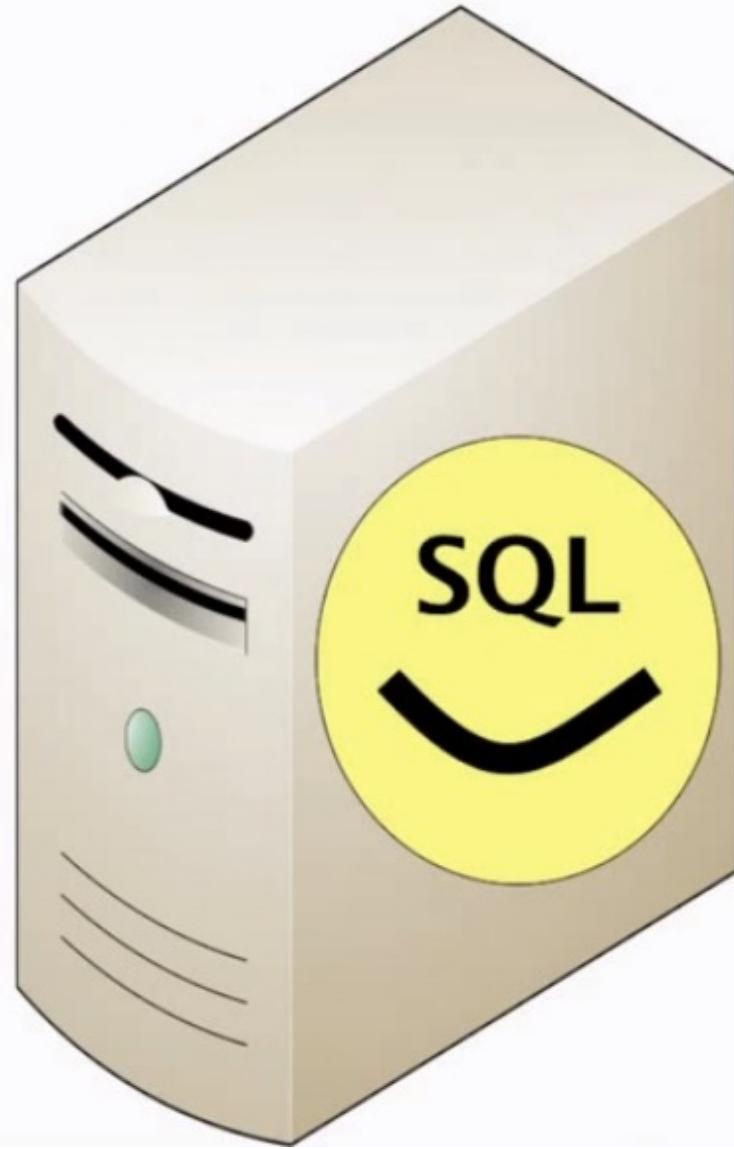
**Lots of
Traffic**





The Google logo, featuring the word "Google" in its signature multi-colored, slightly rounded font.The Amazon.com logo, consisting of the word "amazon.com" in a lowercase, bold, black sans-serif font, with a yellow curved arrow underneath.





Google™



Bigtable

amazon.com®



Dynamo

1980

1990

2000

2010

NoSQL



The Birth of NoSQL

This event has ended!

[View current events hosted by Last.fm](#)

NOSQL meetup

Thursday, June 11, 2009 from 10:00 AM to 5:00 PM (PT)
San Francisco, CA

Ticket Information

TYPE	REMAINING	END	QUANTITY
Free ticket	Sold Out	Ended	Free Sold Out

Share this!



Be the first of your friends to like this.

Event Details

Introduction

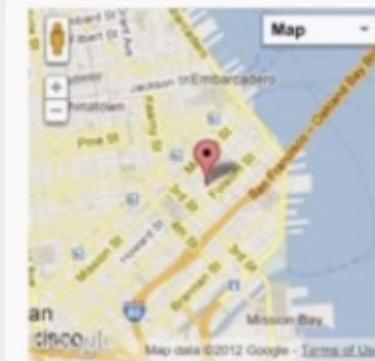
This meetup is about "open source, distributed, non relational databases".

Have you run into limitations with traditional relational databases? Don't mind trading a query language for scalability? Or perhaps you just like shiny new things to try out? Either way this meetup is for you.

Join us in figuring out why these newfangled Dynamo clones and BigTables have become so popular lately. We have gathered presenters from the most interesting projects around to give us all an introduction to the field.

Preliminary schedule

When & Where



CBS interactive, Magma room
235 Second Street
San Francisco, CA 94105

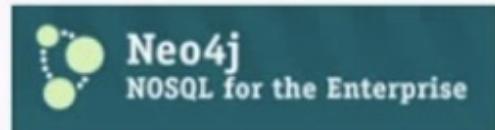
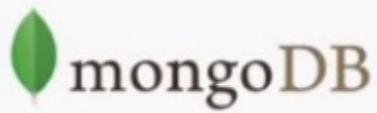
Thursday, June 11, 2009 from 10:00 AM to 5:00 PM (PT)

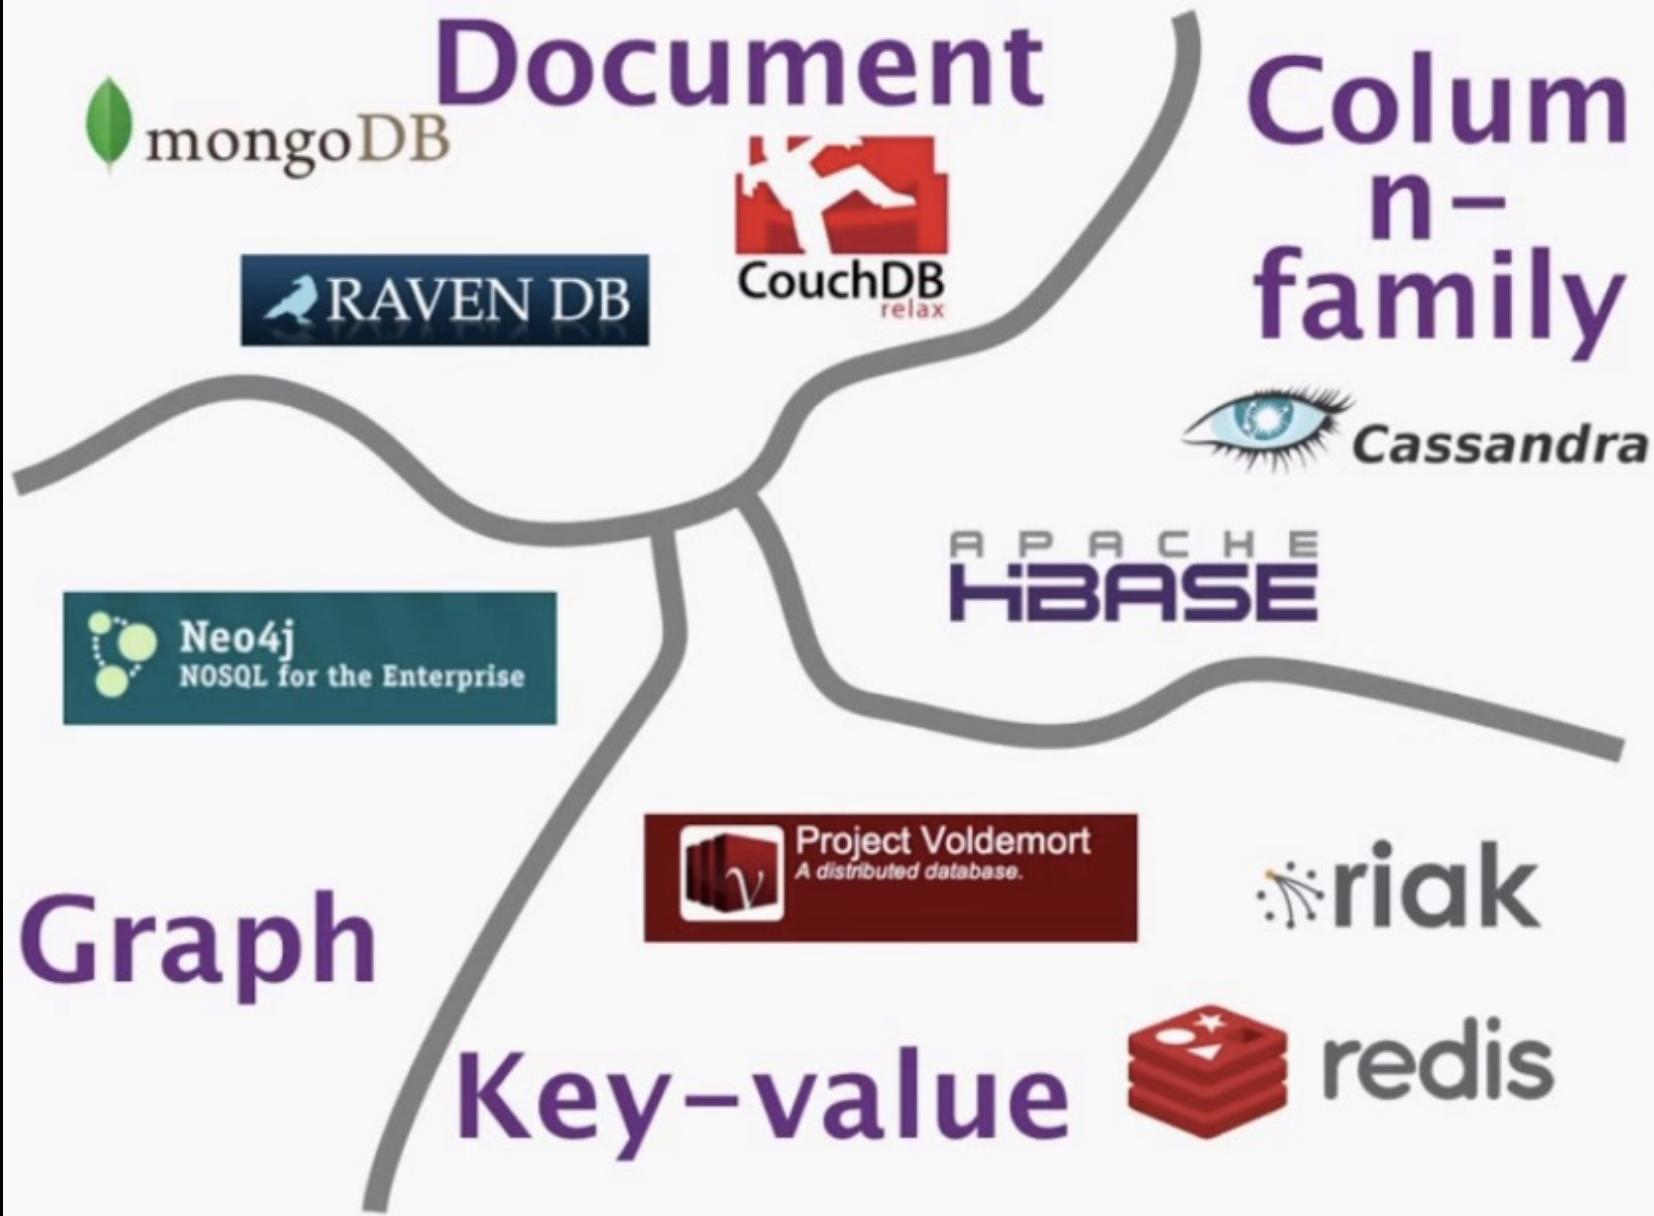
Definition of NoSQL

Characteristics of NoSQL

non-relational open-source
cluster-friendly
21st Century Web
schema-less

Data Model





Key-value

10025



10026



10043



10048

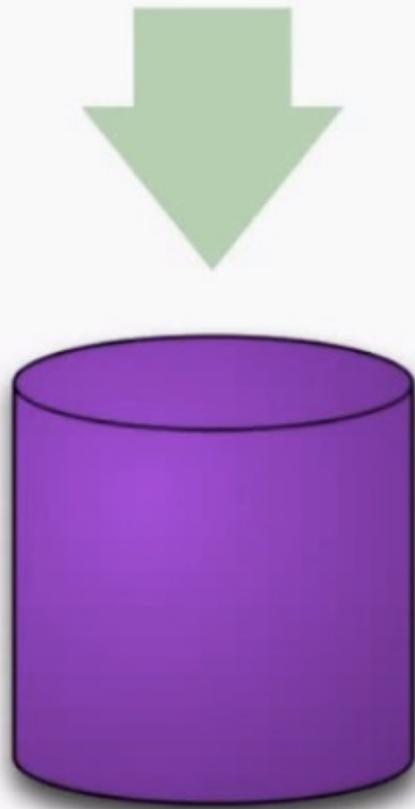


Document

```
{"id": 1001,  
"customer_id": 7231,  
"line-itmes": [  
    {"product_id": 4555, "quantity": 8},  
    {"product_id": 7655, "quantity": 4}, {"product_id": 8755,  
    "quantity": 10}  
],  
"id": 1002,  
"customer_id": 9831,  
"line-itmes": [  
    {"product_id": 4555, "quantity": 3},  
    {"product_id": 2155, "quantity": 4}],  
"discount-code": "Y"}
```

no
schema

`anOrder["price"] * anOrder["quantity"]`

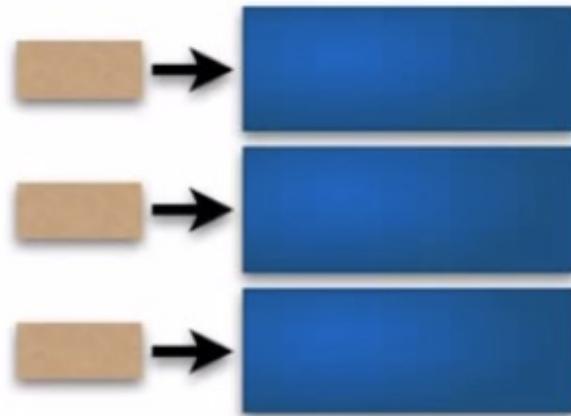


`anOrder["price"] * anOrder["quantity"]`

implicit
schema



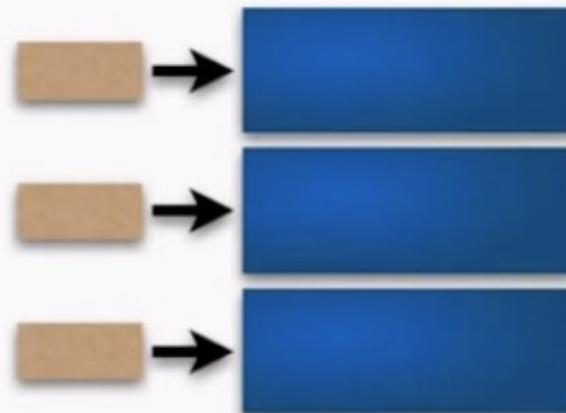
Key-Value



Document

```
{"id": 1001,  
 {"id": 1002,  
 "customer_id": 7231,  
 "line-items": [  
 {"product_id": 4555, "quantity": 8},  
 {"product_id": 7655, "quantity": 4},  
 {"product_id": 8755, "quantity": 3}]  
 }
```

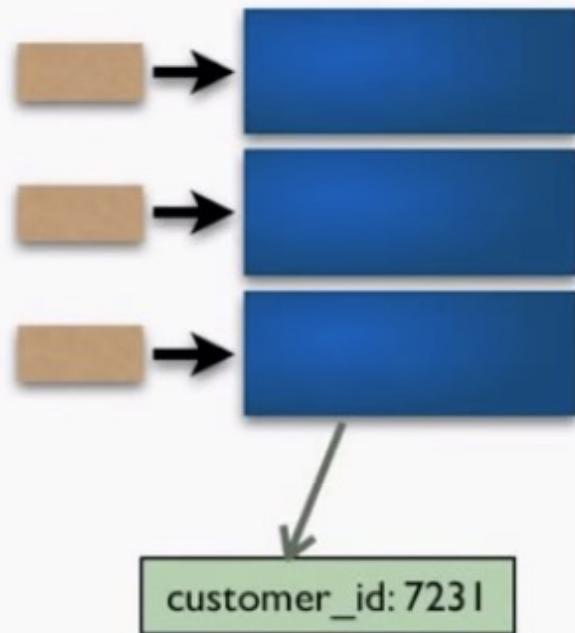
Key-Value



Document

```
{"id": 1001,  
 {"id": 1002,  
 "customer_id": 7231,  
 "line-items": [  
 {"product_id": 4555, "quantity": 8},  
 {"product_id": 7655, "quantity": 4},  
 {"product_id": 8755, "quantity": 3}]  
 }
```

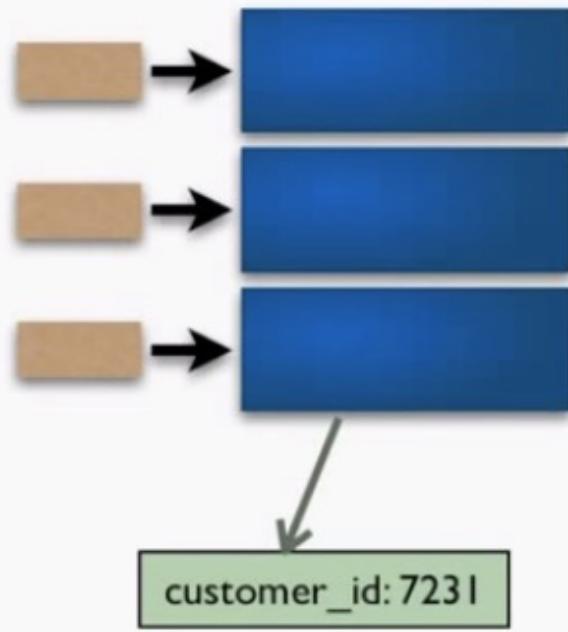
Key-Value



Document

```
{"id": 1001,  
 {"id": 1002,  
 "customer_id": 7231,  
 "line-items": [  
 {"product_id": 4555, "quantity": 8},  
 {"product_id": 7655, "quantity": 4},  
 {"product_id": 8755, "quantity": 3}]  
 }
```

Key-Value

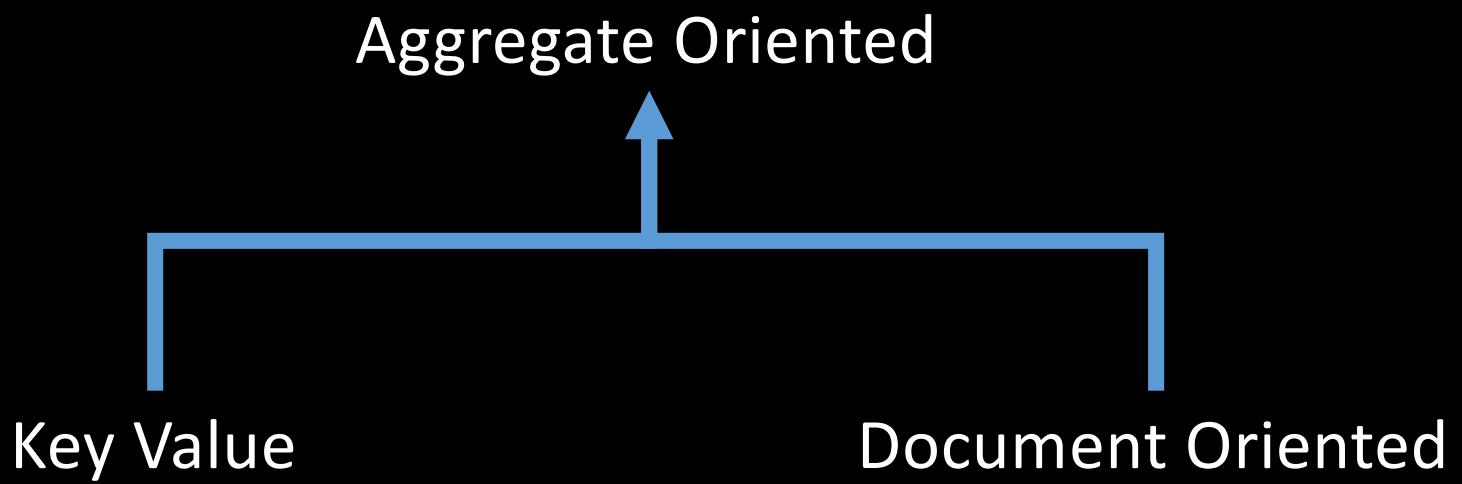


metadata

Document

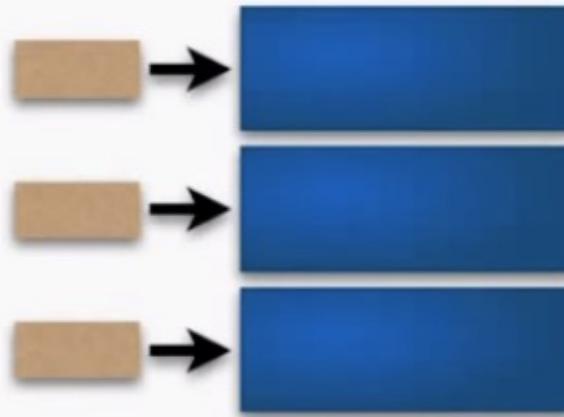
```
{"id": 1001,  
 "id": 1002,  
   "customer_id": 7231,  
   "line-items": [  
     {"product_id": 4555, "quantity": 8},  
     {"product_id": 7655, "quantity": 4},  
     {"product_id": 8755, "quantity": 3}]  
}
```

key





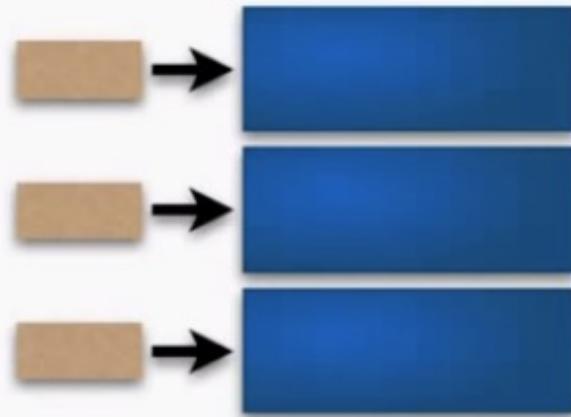
Key-Value



Document

```
{"id": 1001,  
 {"id": 1001,  
 "customer_id": 7231,  
 "line-items": [  
 {"product_id": 4555, "quantity": 8},  
 {"product_id": 7655, "quantity": 4},  
 {"product_id": 8755, "quantity": 3}]}}
```

Key-Value

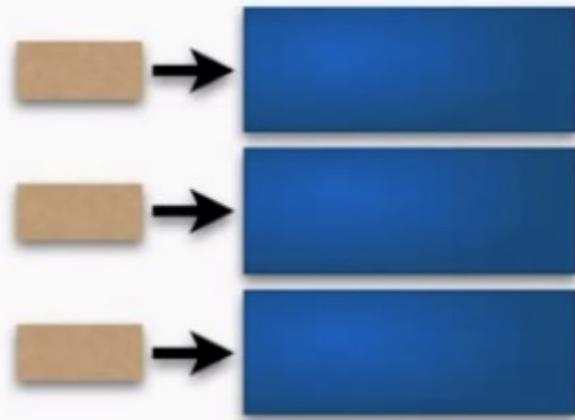


Value ==
Aggregate

Document

```
{"id": 1001,  
 "id": 1001,  
 "customer_id": 7231,  
 "line-items": [  
     {"product_id": 4555, "quantity": 8},  
     {"product_id": 7655, "quantity": 4},  
     {"product_id": 8755, "quantity": 3}]}{
```

Key-Value



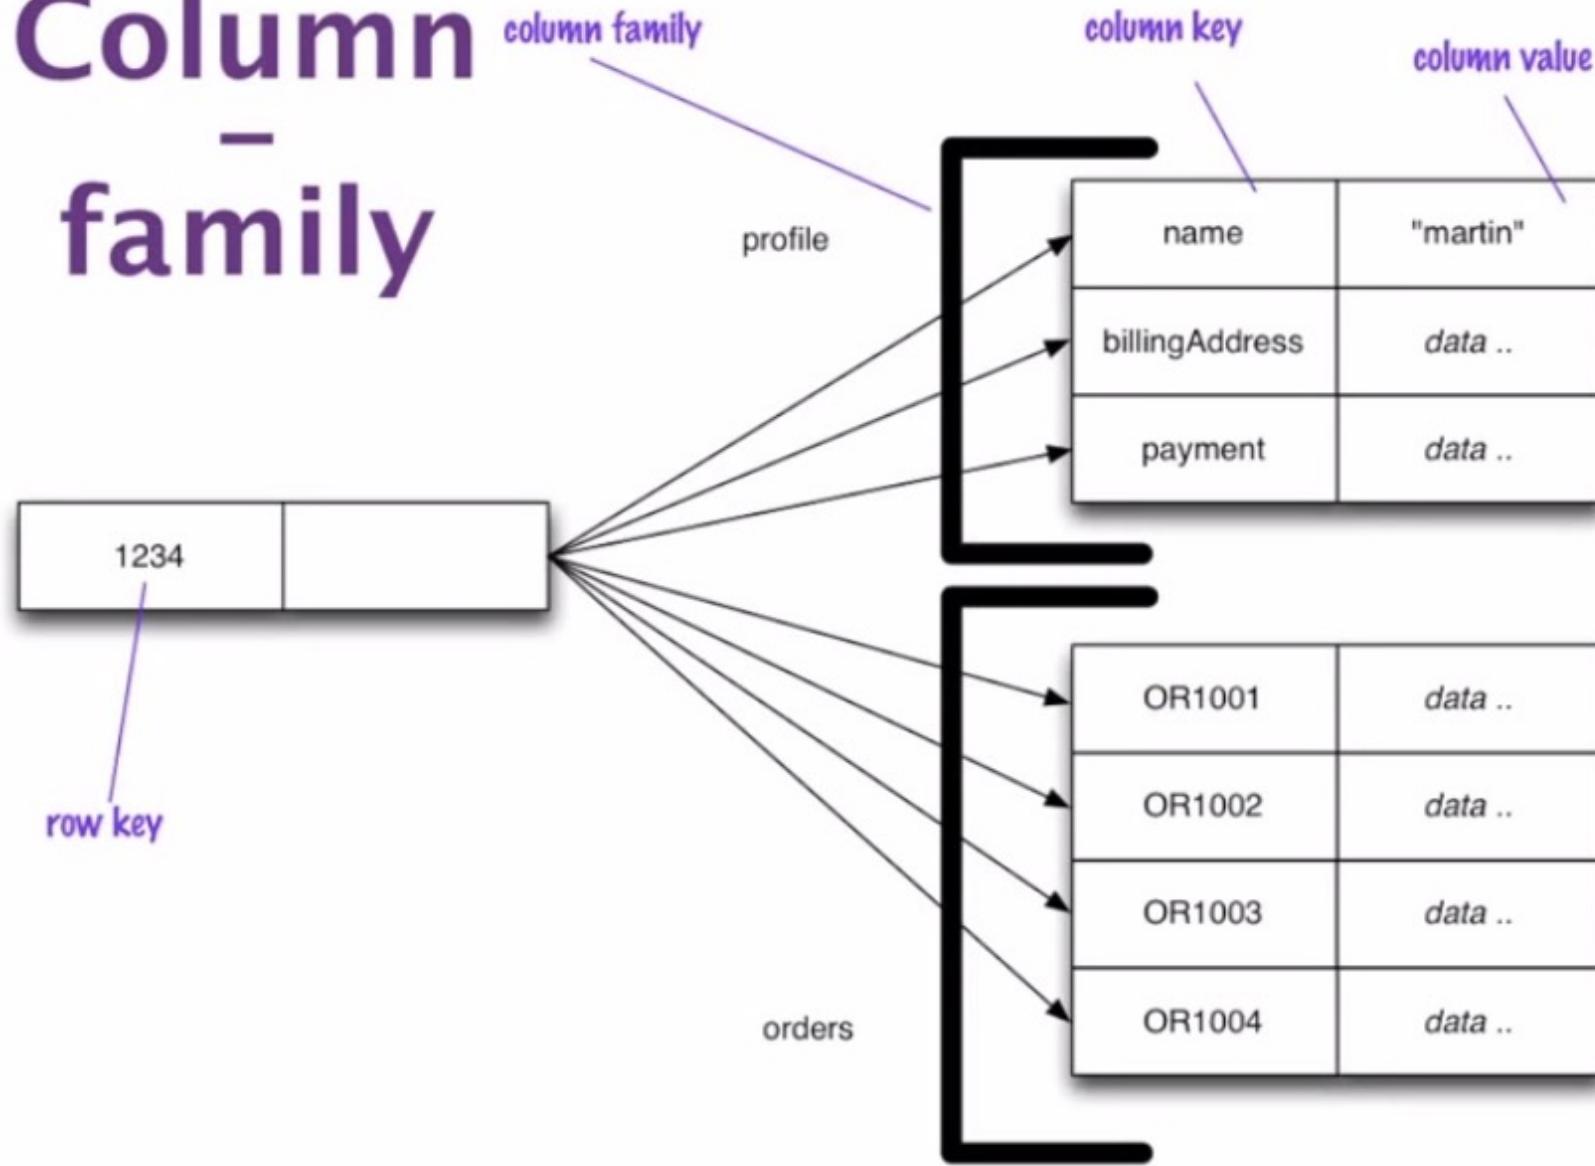
Value ==
Aggregate

Document

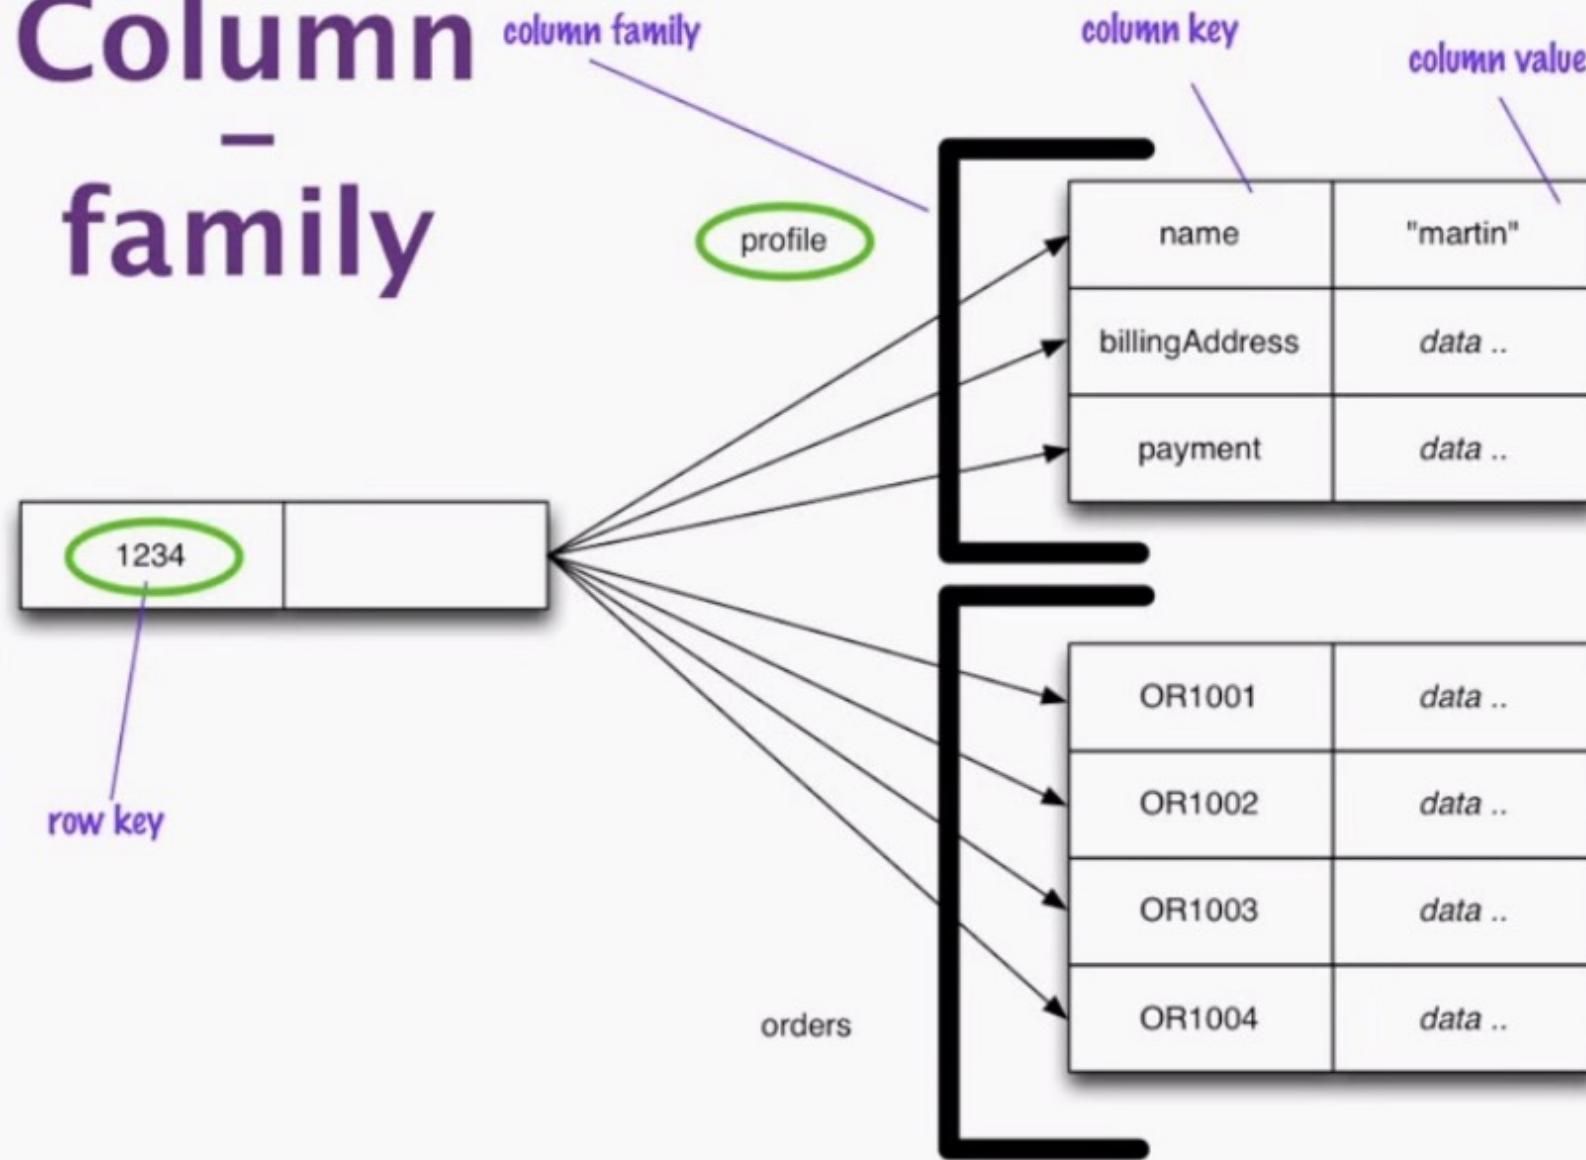
```
{"id": 1001,  
 {"id": 1001,  
 "customer_id": 7231,  
 "line-items": [  
 {"product_id": 4555, "quantity": 8},  
 {"product_id": 7655, "quantity": 4},  
 {"product_id": 8755, "quantity": 3}]}  
}
```

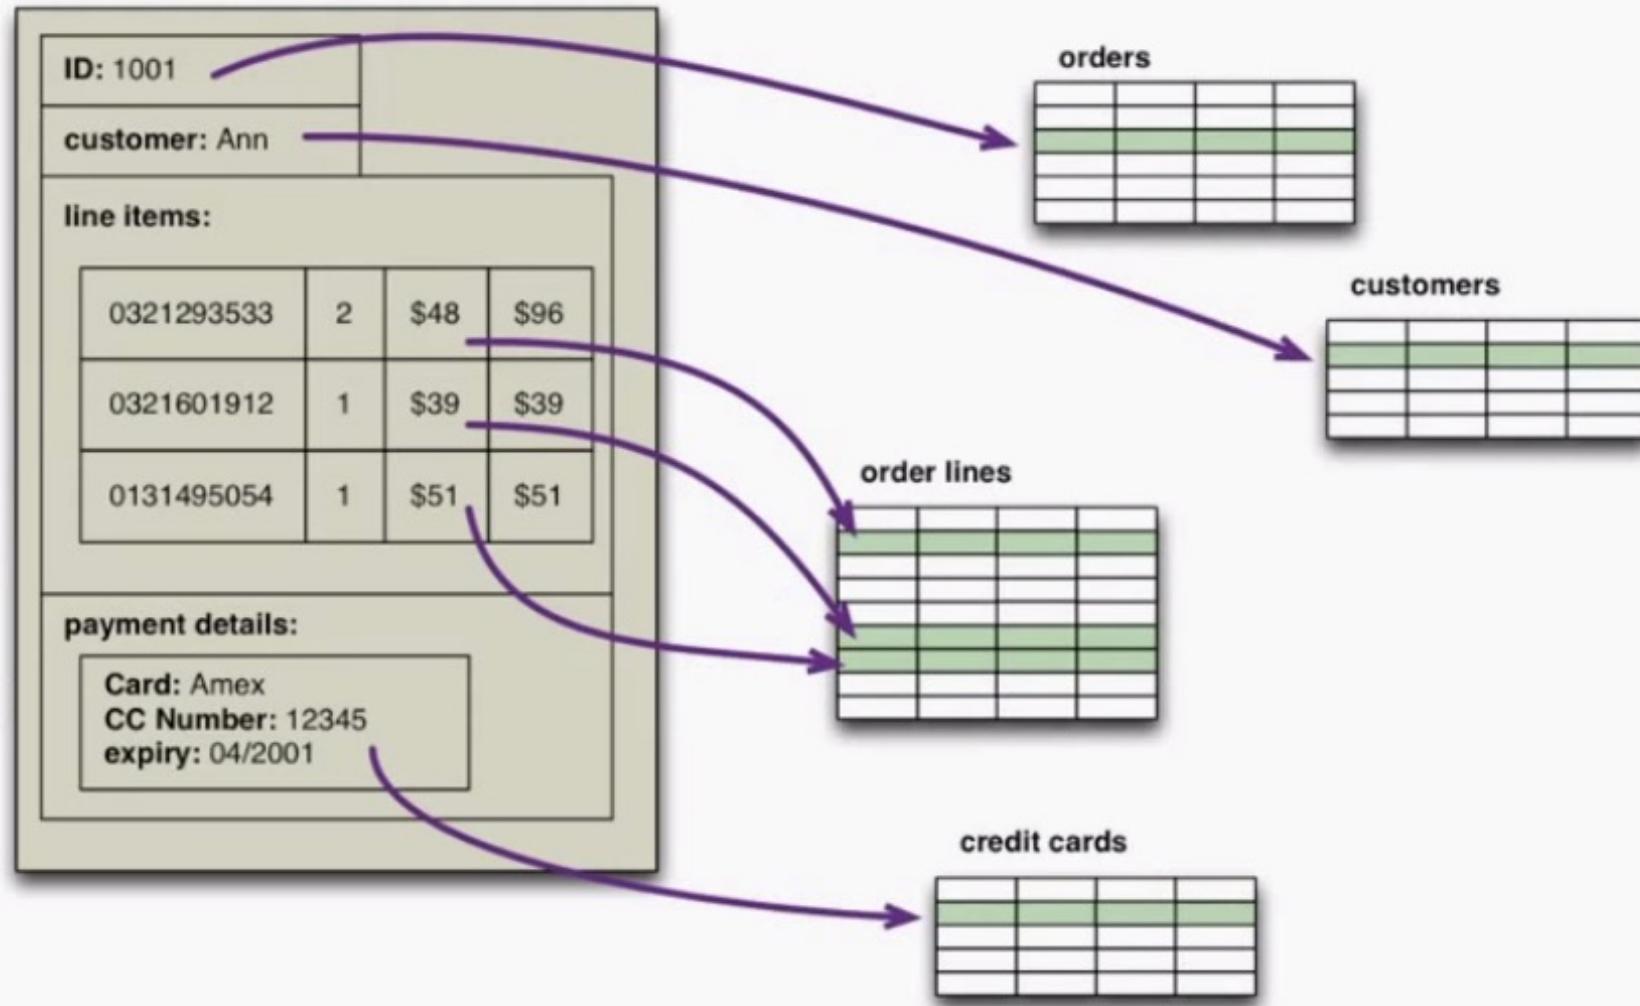
Document ==
Aggregate

Column - family

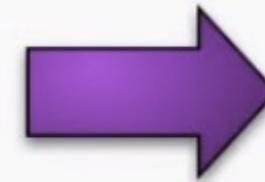


Column - family



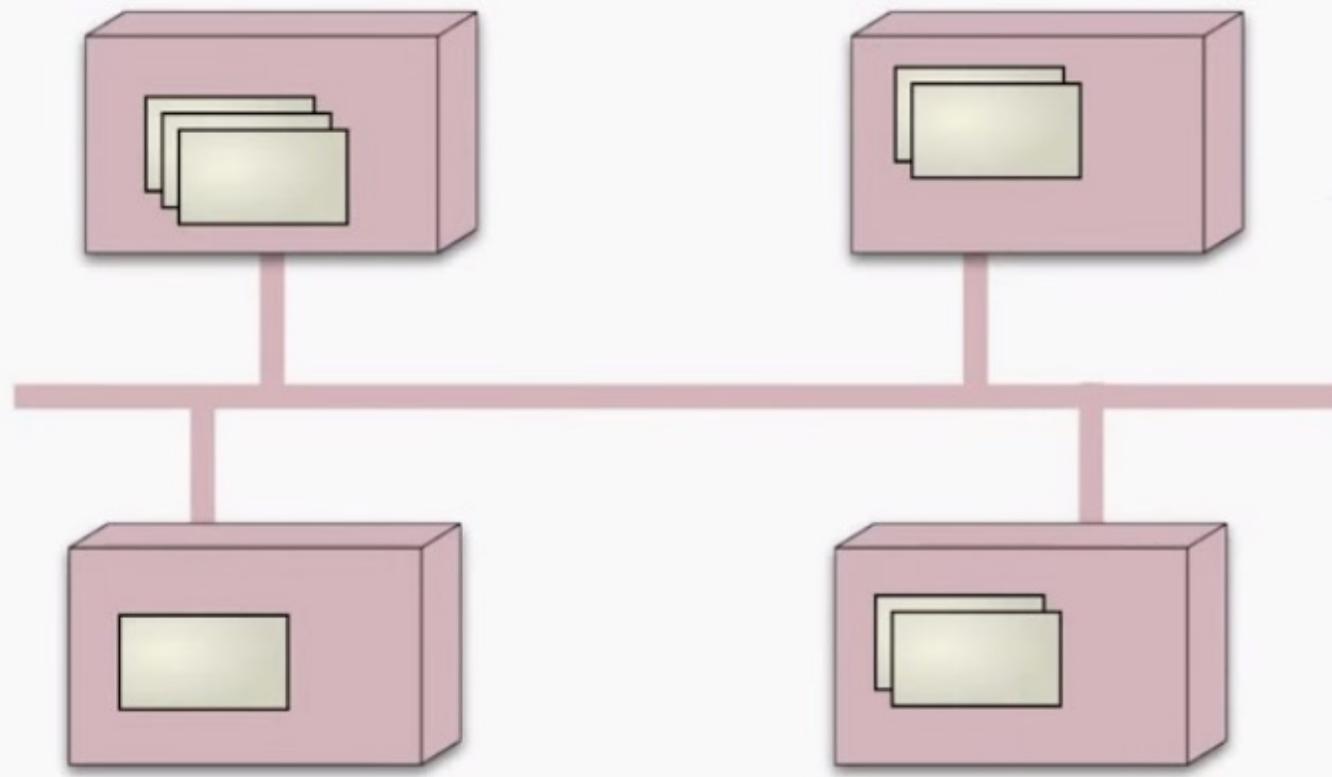


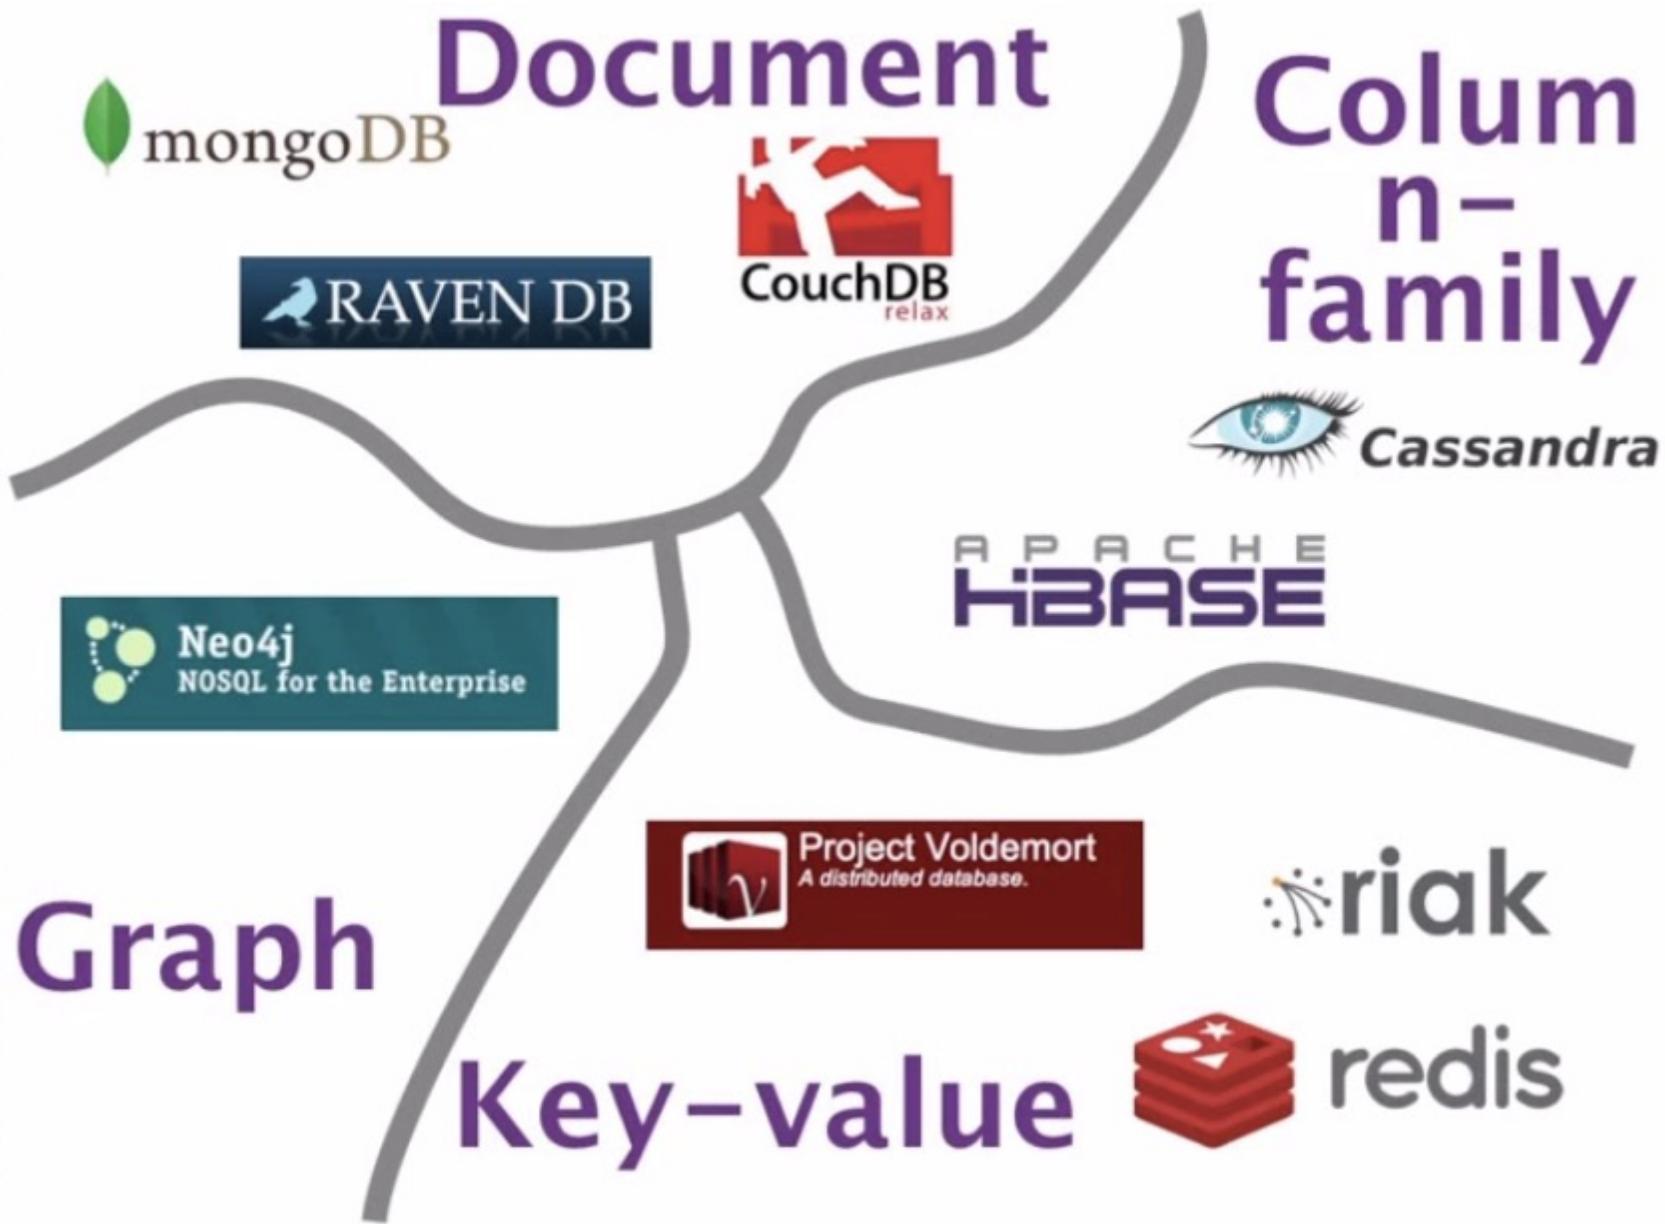
ID: 1001												
customer: Ann												
line items:												
<table border="1"><tr><td>0321293533</td><td>2</td><td>\$48</td><td>\$96</td></tr><tr><td>0321601912</td><td>1</td><td>\$39</td><td>\$39</td></tr><tr><td>0131495054</td><td>1</td><td>\$51</td><td>\$51</td></tr></table>	0321293533	2	\$48	\$96	0321601912	1	\$39	\$39	0131495054	1	\$51	\$51
0321293533	2	\$48	\$96									
0321601912	1	\$39	\$39									
0131495054	1	\$51	\$51									
payment details:												
Card: Amex CC Number: 12345 expiry: 04/2001												



aggregate







Aggregate-
Oriented

Document
Key-value

Column-
family

Graph

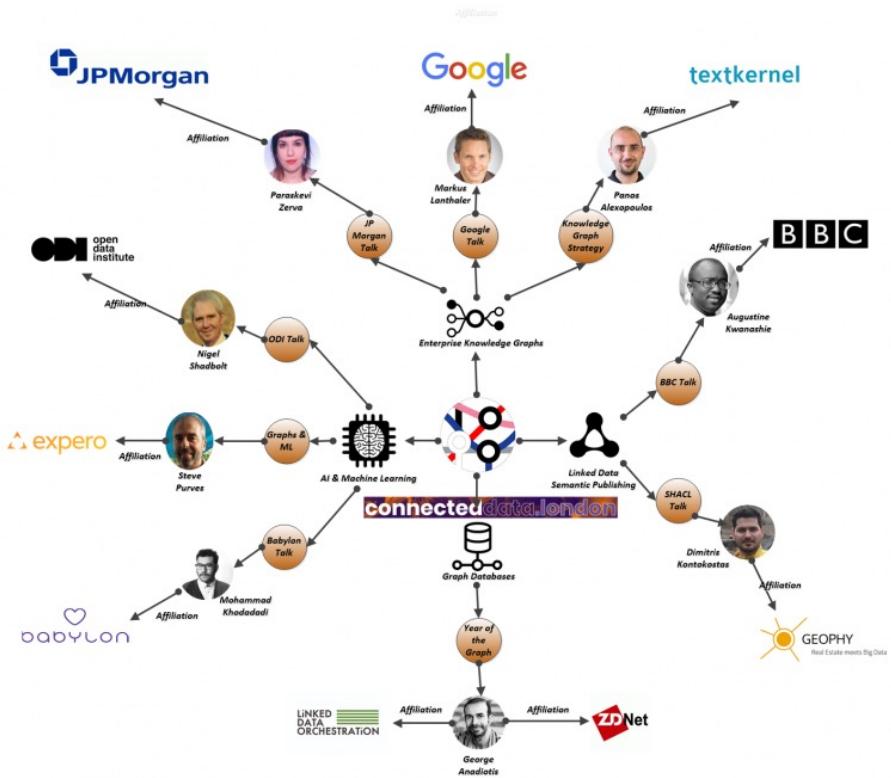
Humans think in graphs

- We understand better things by matching them to things that we know.
- And the things we know ... are in the form of graphs



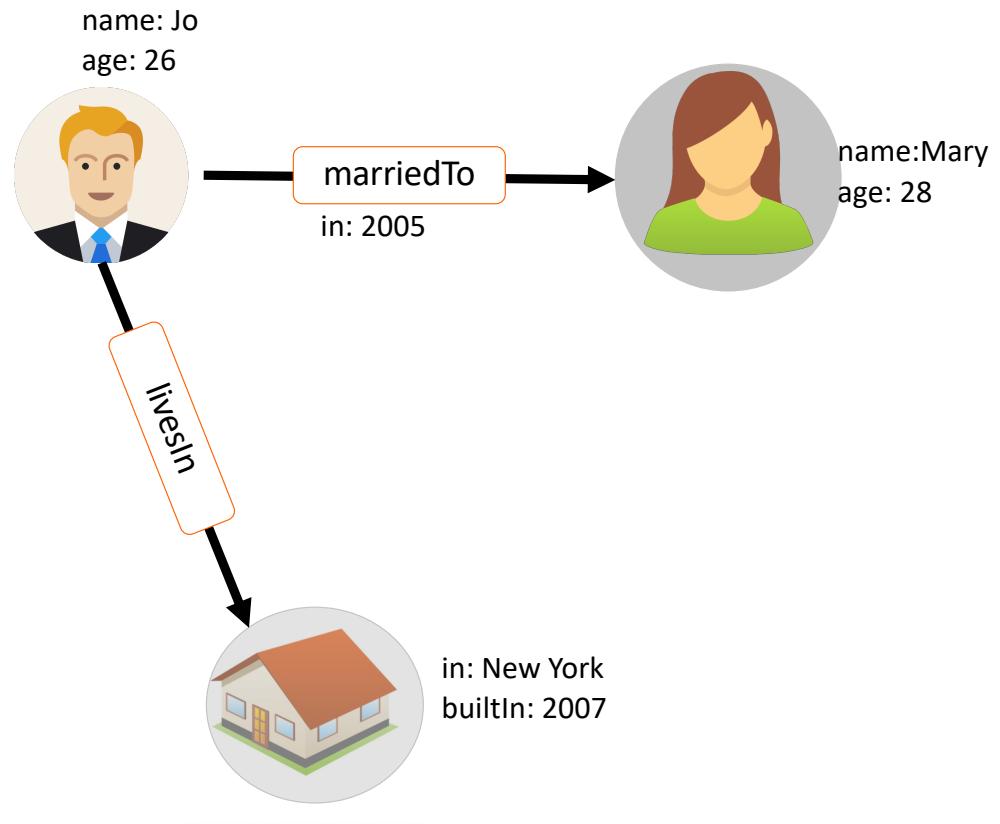
Graphs capture naturally heterogeneous data

- No specific schema
- Shows how different parts connect to each other.
- Hard to model with relational joins
- Very important for more informative insights



The property Graph Model

- Nodes
 - Represent Objects
 - Can be labeled
- Relationships
 - Relate Nodes
 - Have a label (used as a type)
 - Have direction
- Properties
 - <Name: value> pairs
 - Describe characteristics
 - Can be on nodes or on edges



Graph Data: Social Networks



Facebook social graph

4-degrees of separation [Backstrom-Boldi-Rosa-Ugander-Vigna, 2011]

Gremlin

- Java, Rest, Cypher, Gremlin
- Gremlin: Graph traversal language
- g: variable that represents a graph
- Profiles

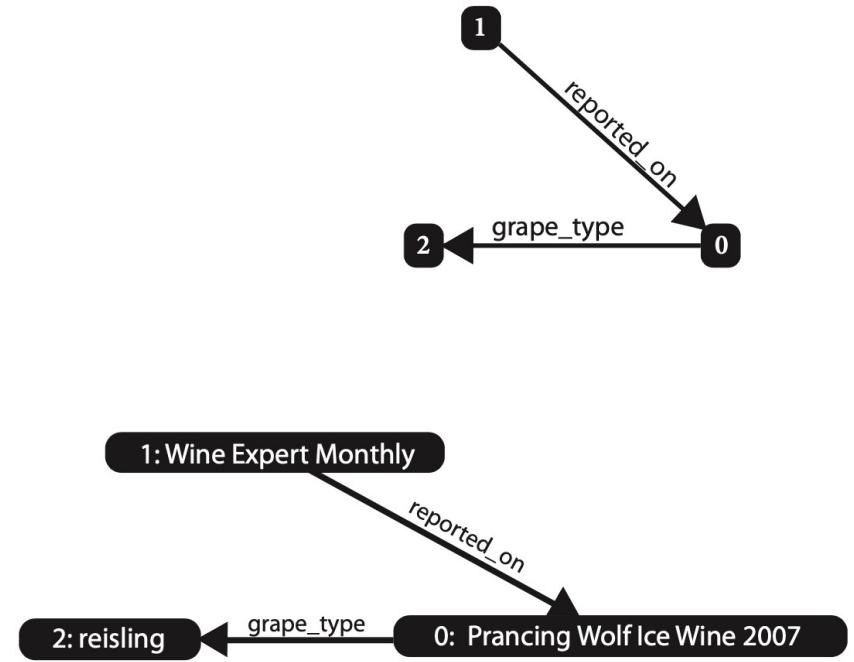
```
gremlin> g.V  
==>v[0]  
==>v[1]  
==>v[2]
```

```
gremlin> g.E  
==> e[0][1-reported_on->0]  
==> e[1][0-grape_type->2]
```

```
gremlin> g.v(0)  
==> v[0]
```

List properties:

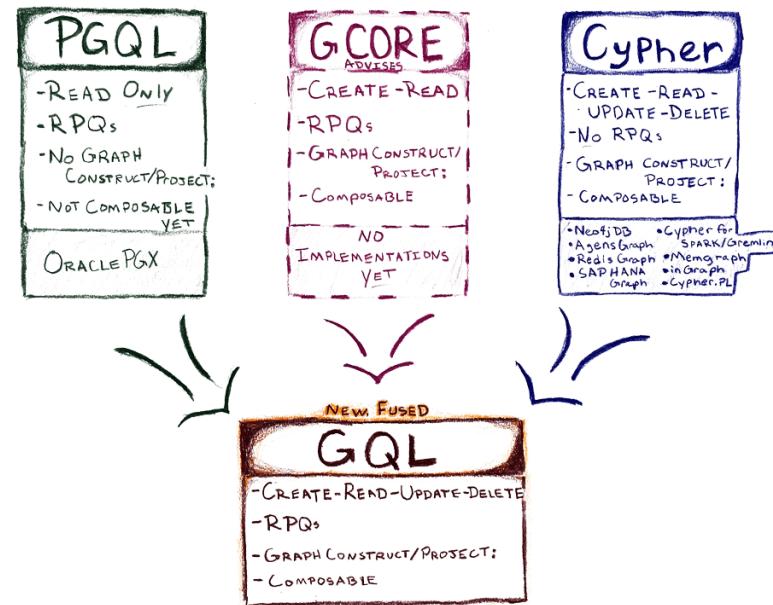
```
gremlin> g.v(0).map()  
==> name=Prancing Wolf Ice Wine 2007
```



The Property Graph Query Language

- Why not SQL
 - Cannot express graph patterns
 - Not easy to do recursions
 - Go easier with heterogeneous structures
- Conceptually easier to understand

```
...
MATCH (a)
CALL {
    MATCH (a) - [:FRIEND_OF] - (x) - [:FRIEND_OF] - (b)
    WHERE a <> x AND x <> b AND a <> b
    RETURN count(DISTINCT b) AS num_foofs
}
RETURN a.name, num_foofs
ORDER BY num_foofs DESC
LIMIT 100
```



Cypher

What Breweries Near Me That My Friends Also Like Have The Highest Rating?



```
1 MATCH (c:Category)<-[ :IN_CATEGORY ]-(brew:Business)
2 MATCH (brew)<-[ :REVIEWS ]-(r:Review)<-[ :WROTE ]-(u:User)-[:FRIENDS]-(me:User {user_id: 'qZB8BZ2ZzMF0VfIGHttu7w'})
3 WHERE c.name CONTAINS 'Breweries' AND brew.city = "Phoenix" AND r.stars > 4
4 MATCH (brew)<-[ :REVIEWS ]-(r2:Review)
5 RETURN brew.name, brew.address, avg(r2.stars) AS average_reviews ORDER BY average_reviews DESC
```

- MATCH ...
- (node) - [:RELATIONSHIP] -> (node)
- Predicates
- Aggregations
- Ordering
- Graph vs tabular data

Graph Management Systems



Visit <https://people.cs.aau.dk/~matteo/gdb.html> for a comparison

Aggregate-Oriented

Document

Column-family

Key-value

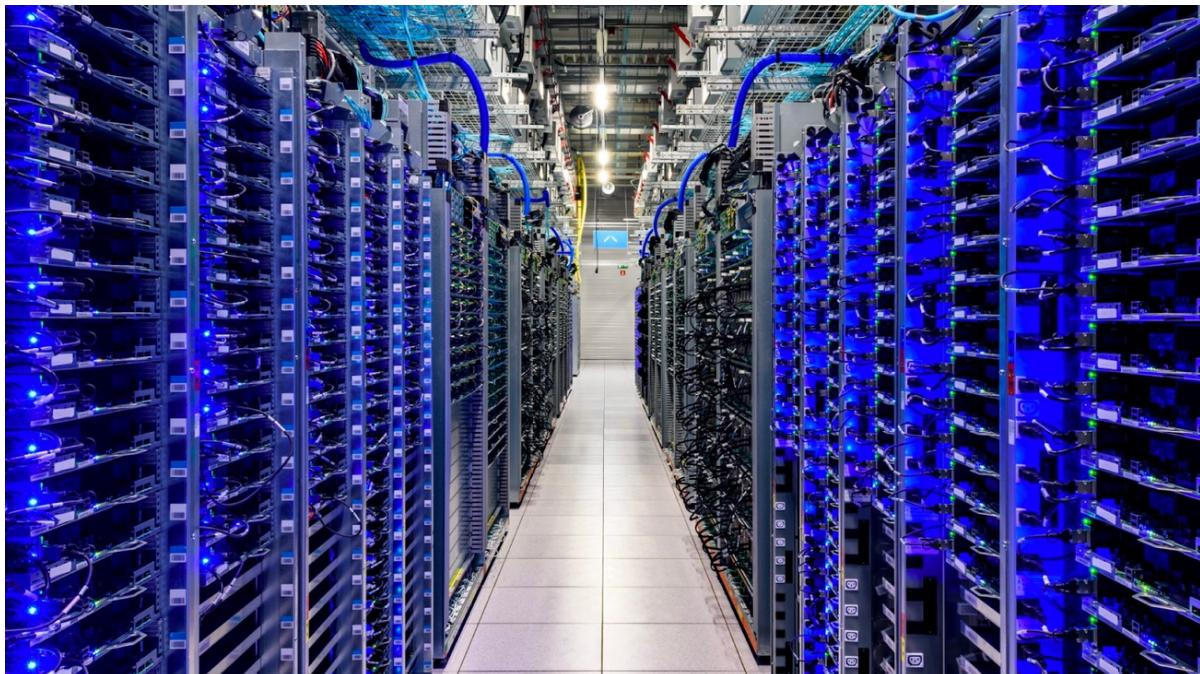
Graph

Schemaless



A file system for Clusters

A Data Center



A Container style Data Center



A google data center may contain 45 containers which is ~60,000 servers





The Need for Consistency ... or Not

ACID

- ACID
 - **Atomic** – Each transaction is either properly carried out or the process halts and the database reverts back to the state before the transaction started. This ensures that all data in the database is valid.
 - **Consistent** – A processed transaction will never endanger the structural integrity of the database.
 - **Isolated** – Transactions cannot compromise the integrity of other transactions by interacting with them while they are still in progress.
 - **Durable** – The data related to the completed transaction will persist even in the cases of network or power outages. If a transaction fails, it will not impact the manipulated data.

BASE

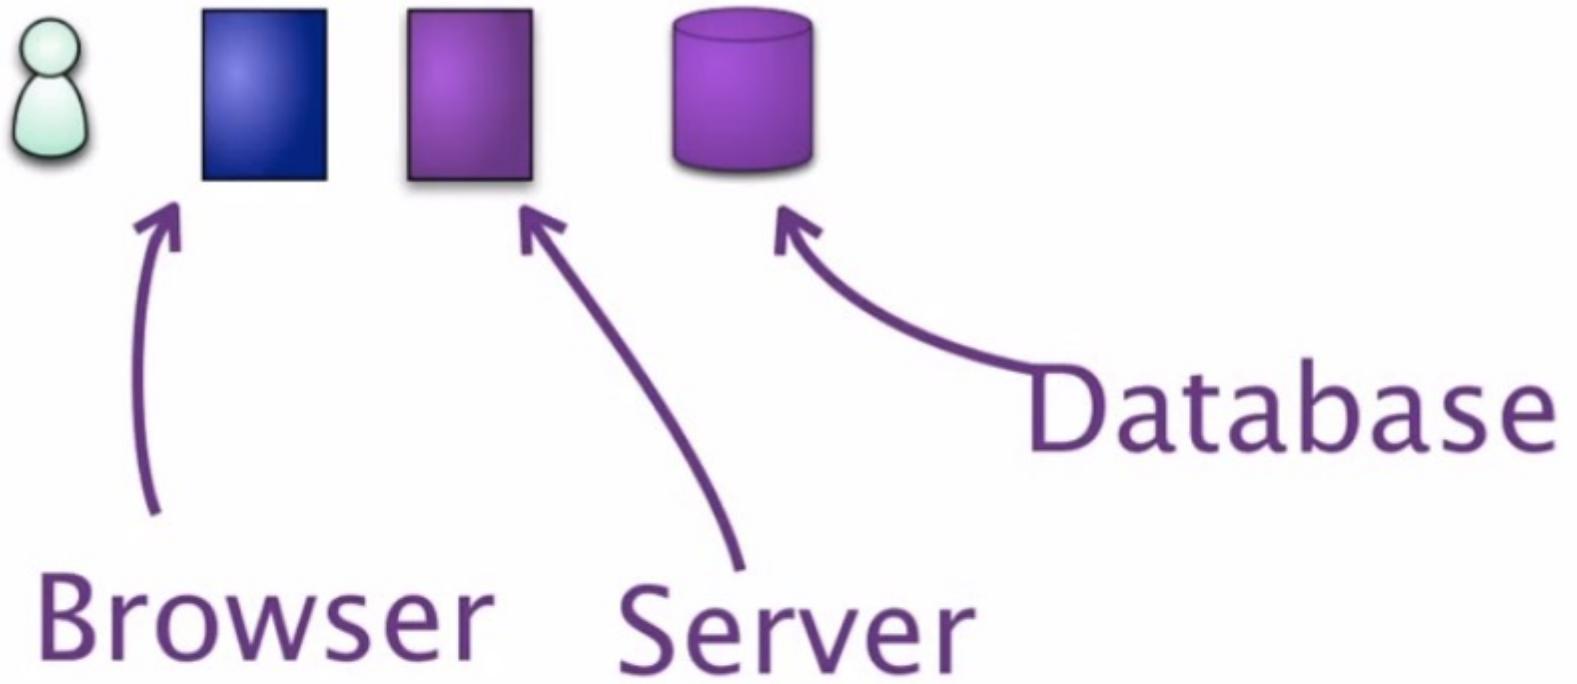
- Base
 - **Basically Available** – Rather than enforcing immediate consistency, BASE-modelled NoSQL databases will ensure availability of data by spreading and replicating it across the nodes of the database cluster.
 - **Soft State** – Due to the lack of immediate consistency, data values may change over time. The BASE model breaks off with the concept of a database which enforces its own consistency, delegating that responsibility to developers.
 - **Eventually Consistent** – The fact that BASE does not enforce immediate consistency does not mean that it never achieves it. However, until it does, data reads are still possible (even though they might not reflect the reality).

BASE

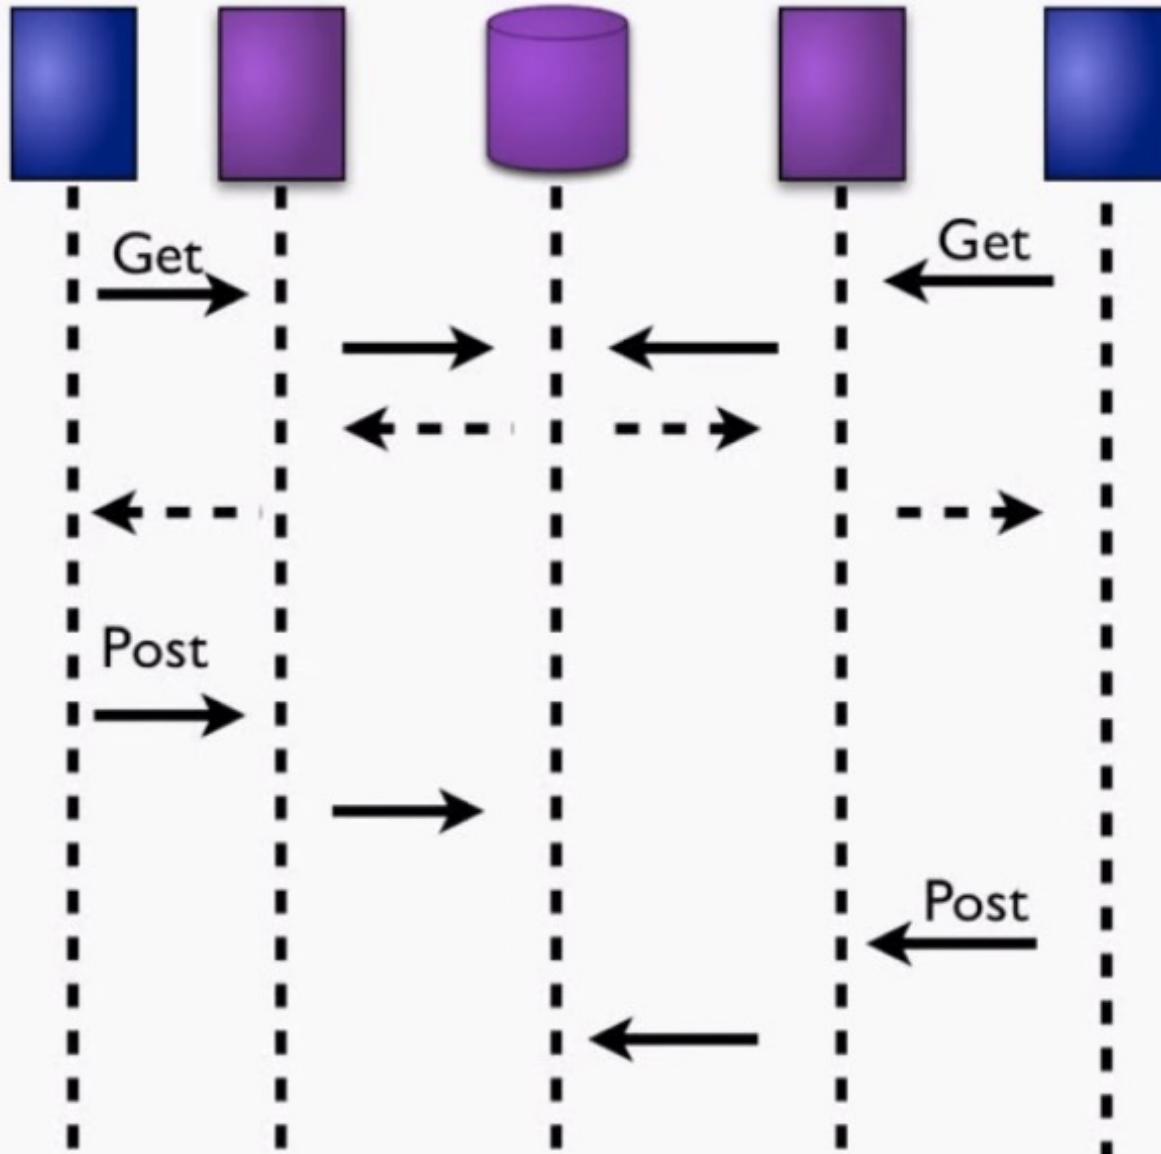
- RDBMS = ACID
- NoSQL = BASE (Well almost)

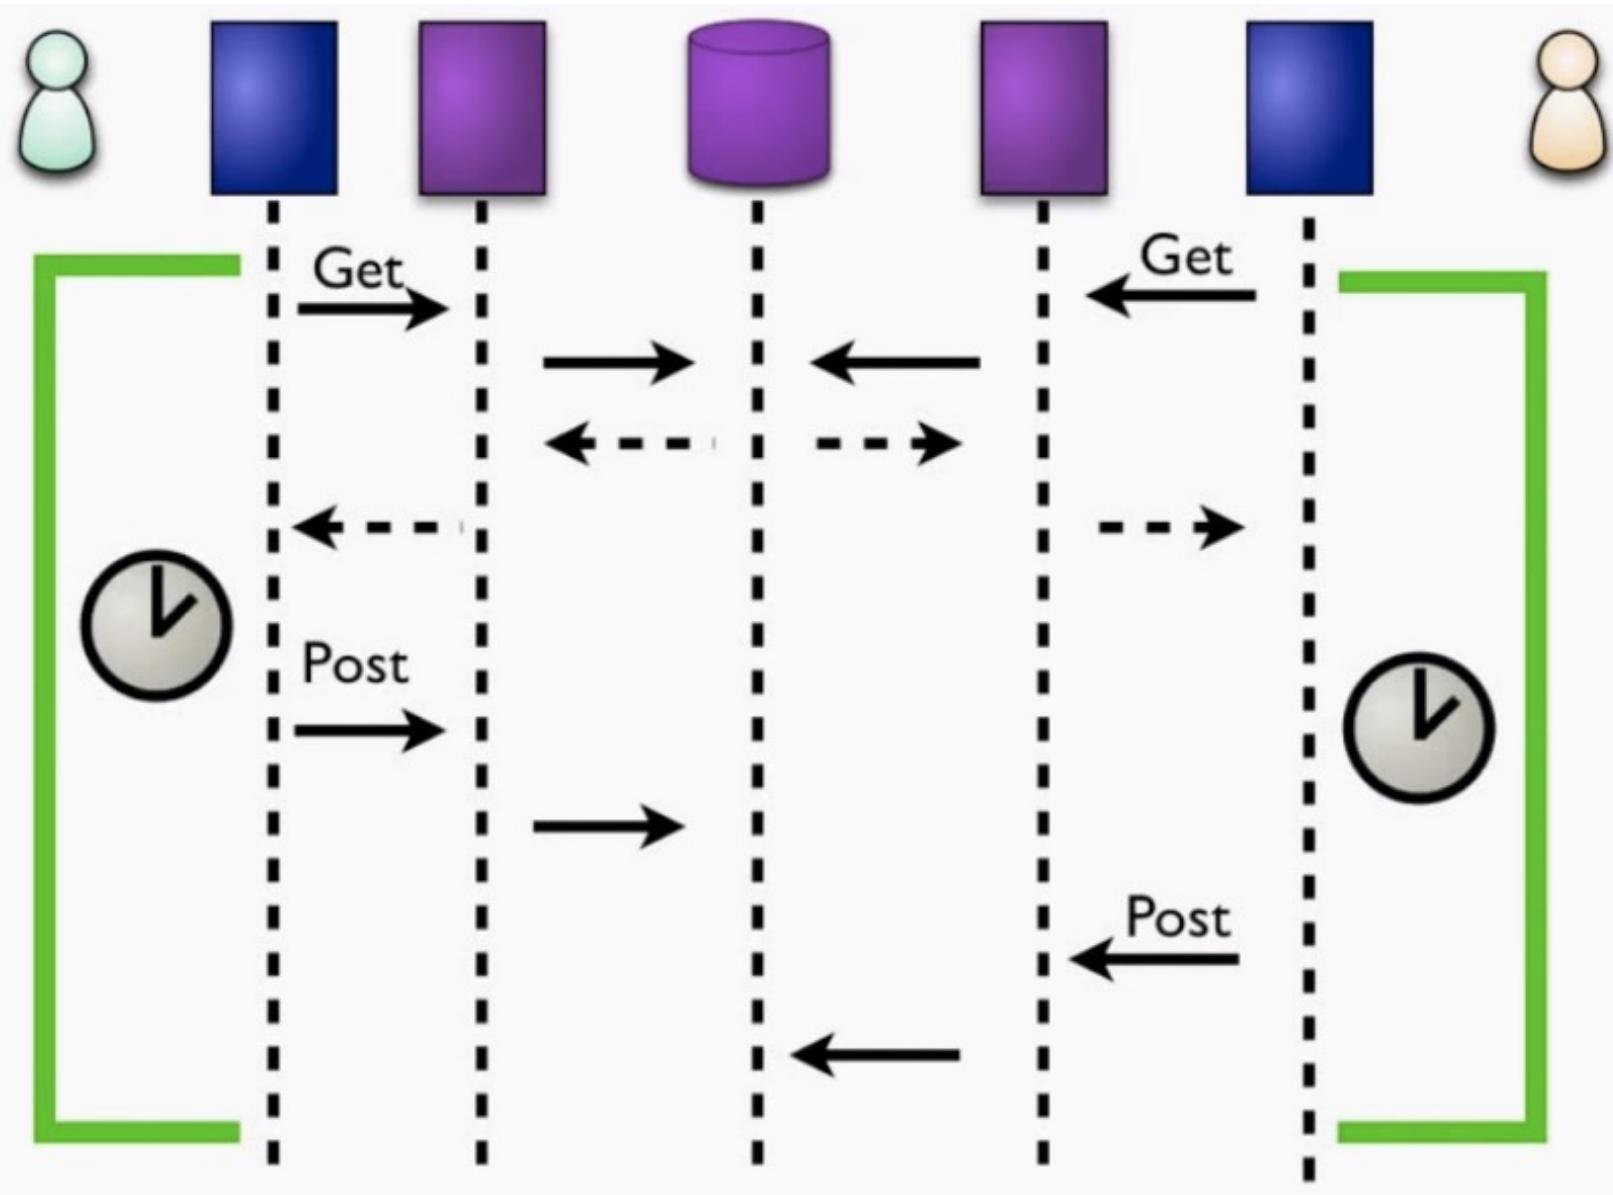
**Aggregate-
Oriented
Document
Column-
family
Key-value**

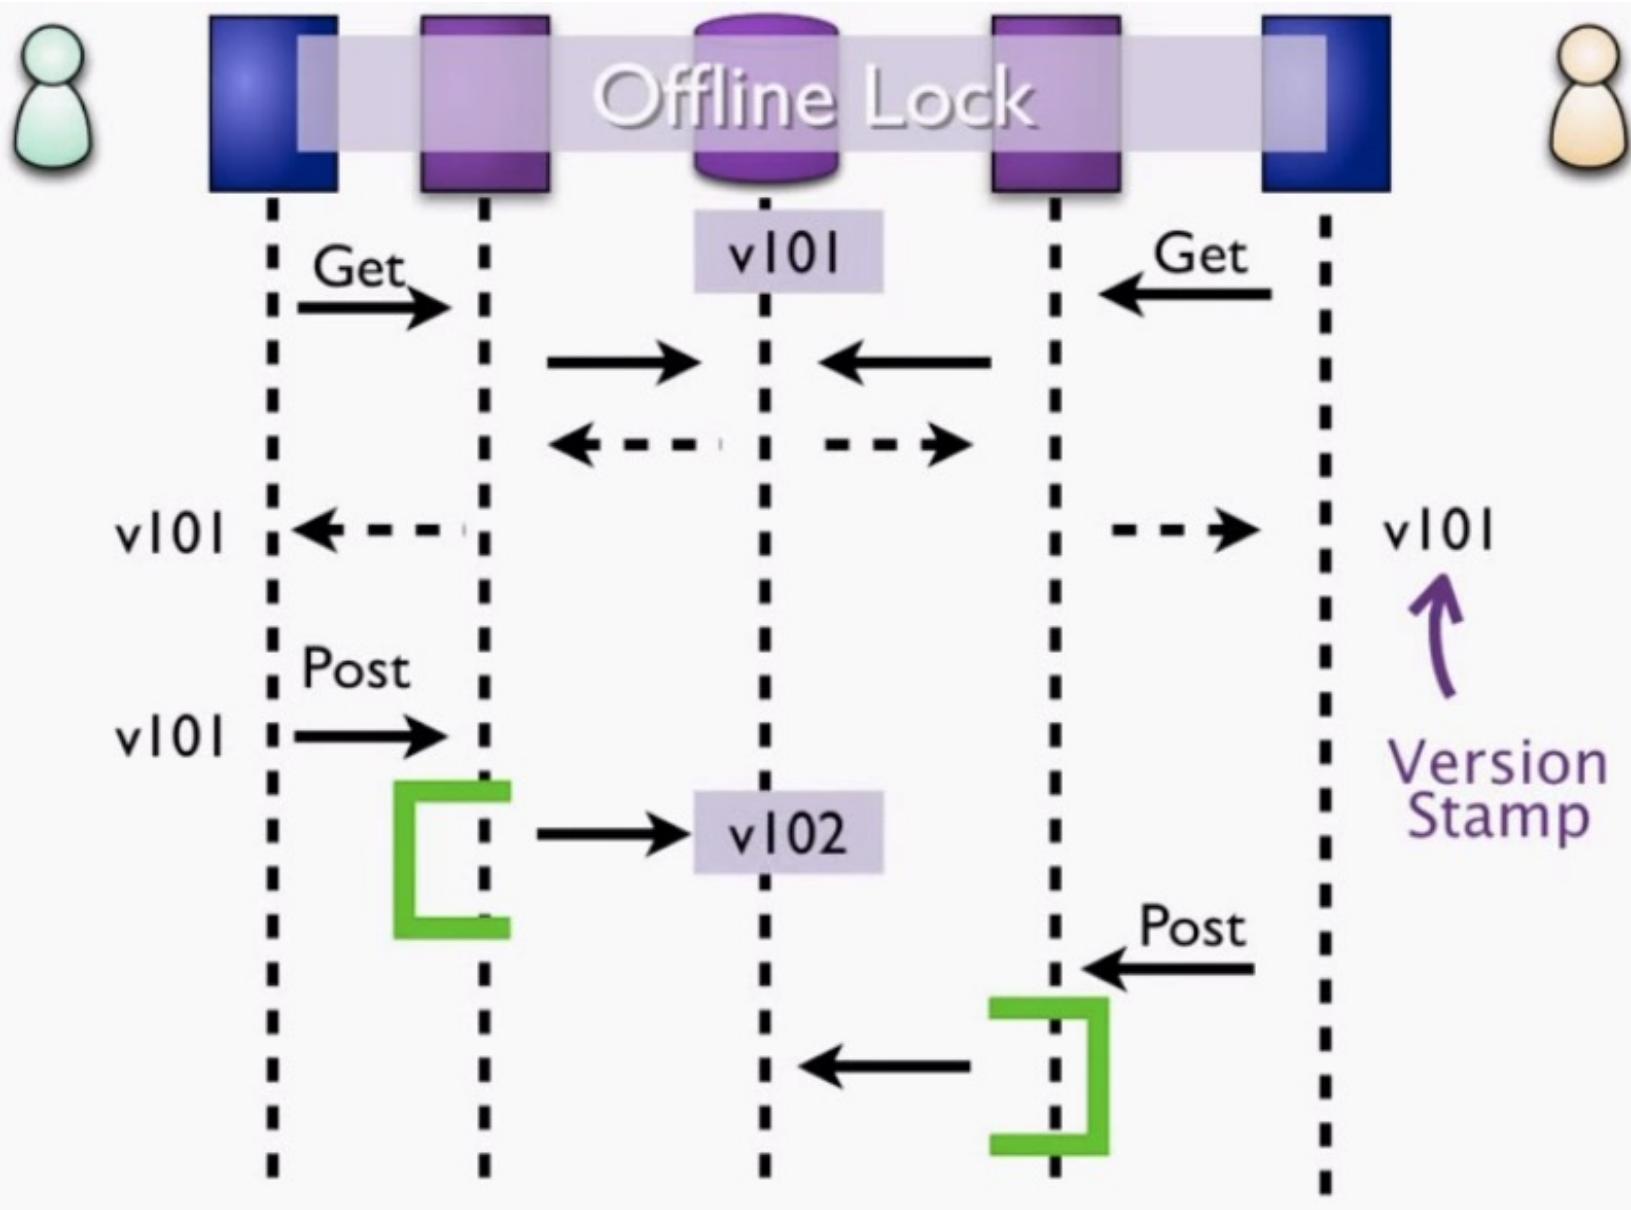
**ACID
Graph**











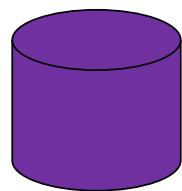


The Effect of Replication

The Effect of Replication

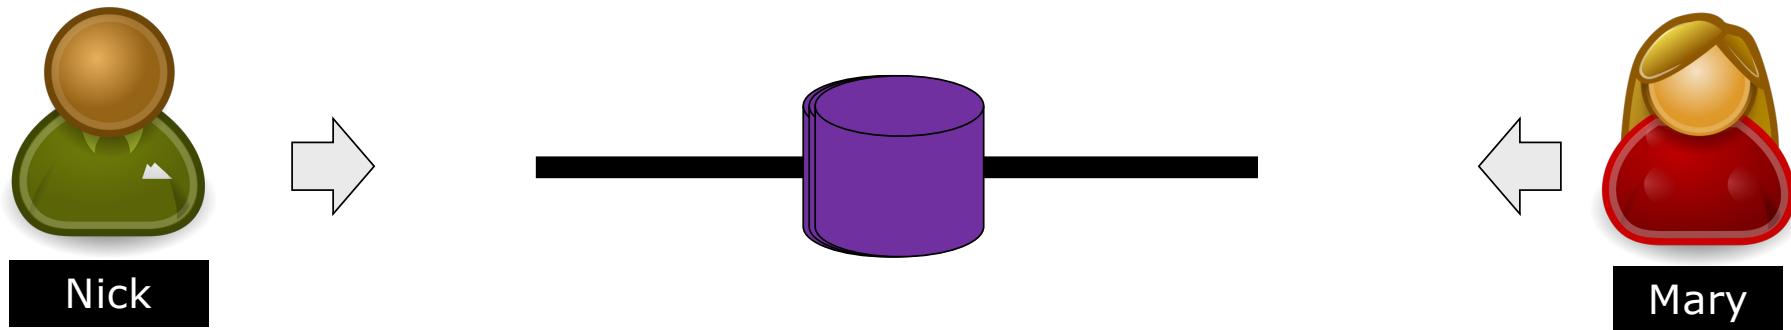


Nick

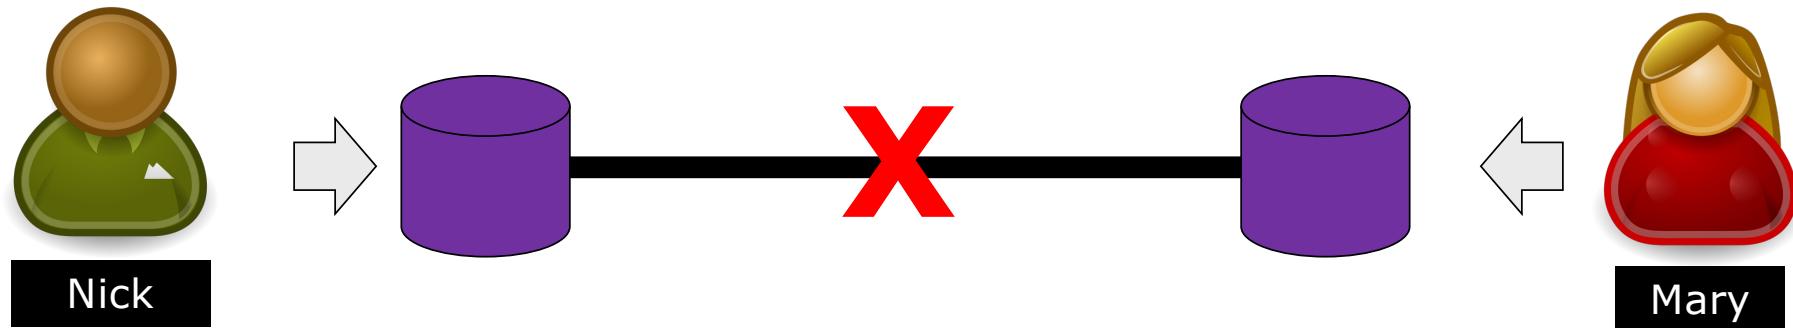


Mary

The Effect of Replication



The Effect of Replication

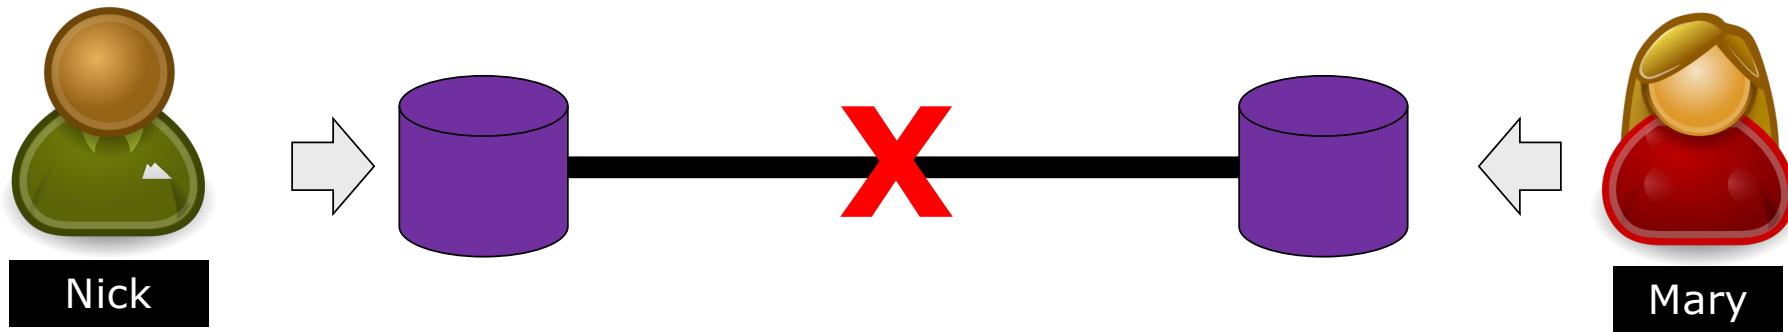


The Effect of Replication



Consistency Guaranteed

The Effect of Replication



Availability Guaranteed

- **Consistency**
 - Do all applications see all the same data?
- **Availability**
 - If some nodes fail, does everything still work?
- **Partitioning**
 - If your nodes can't talk to each other, does everything still work?

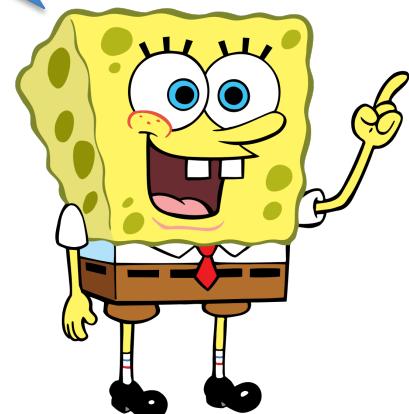
Partition Tolerance

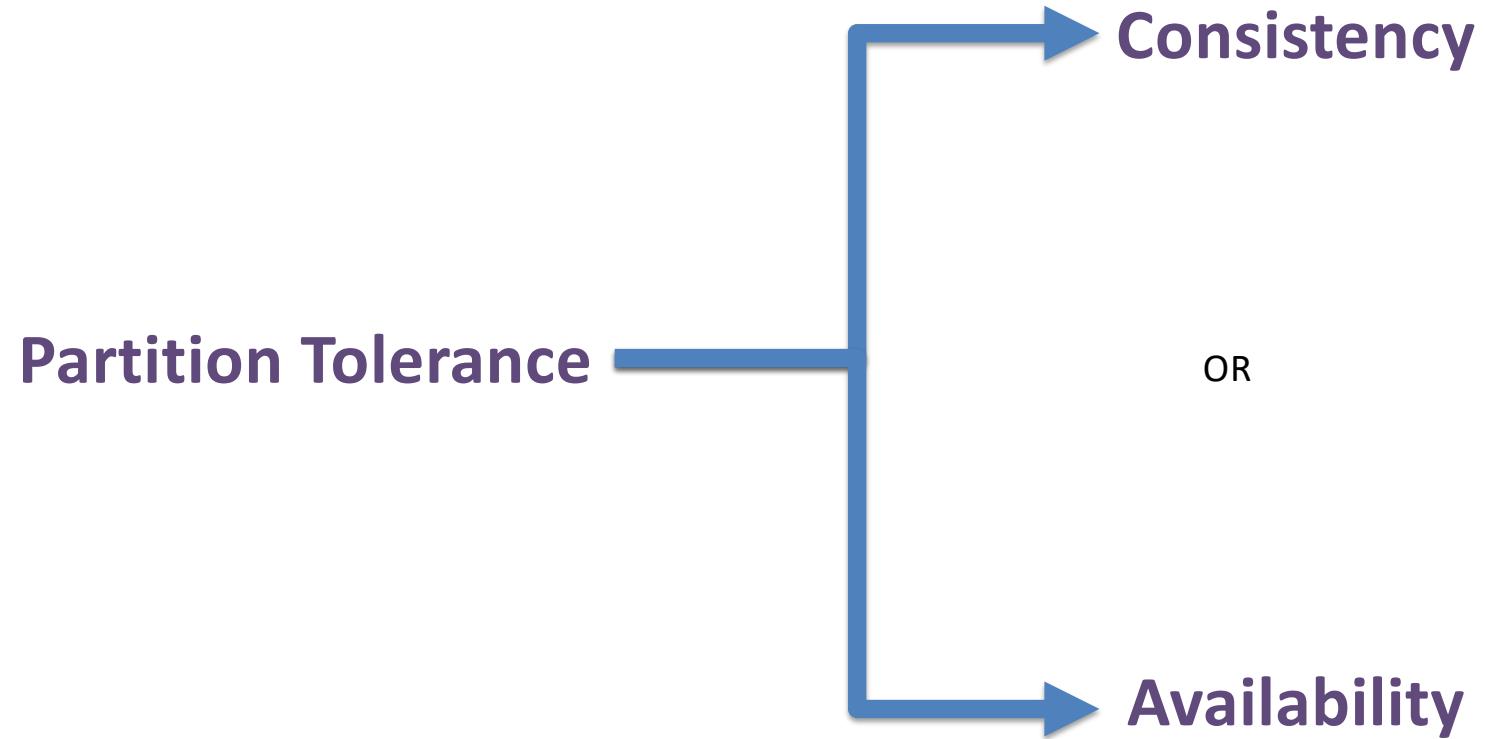
Consistency

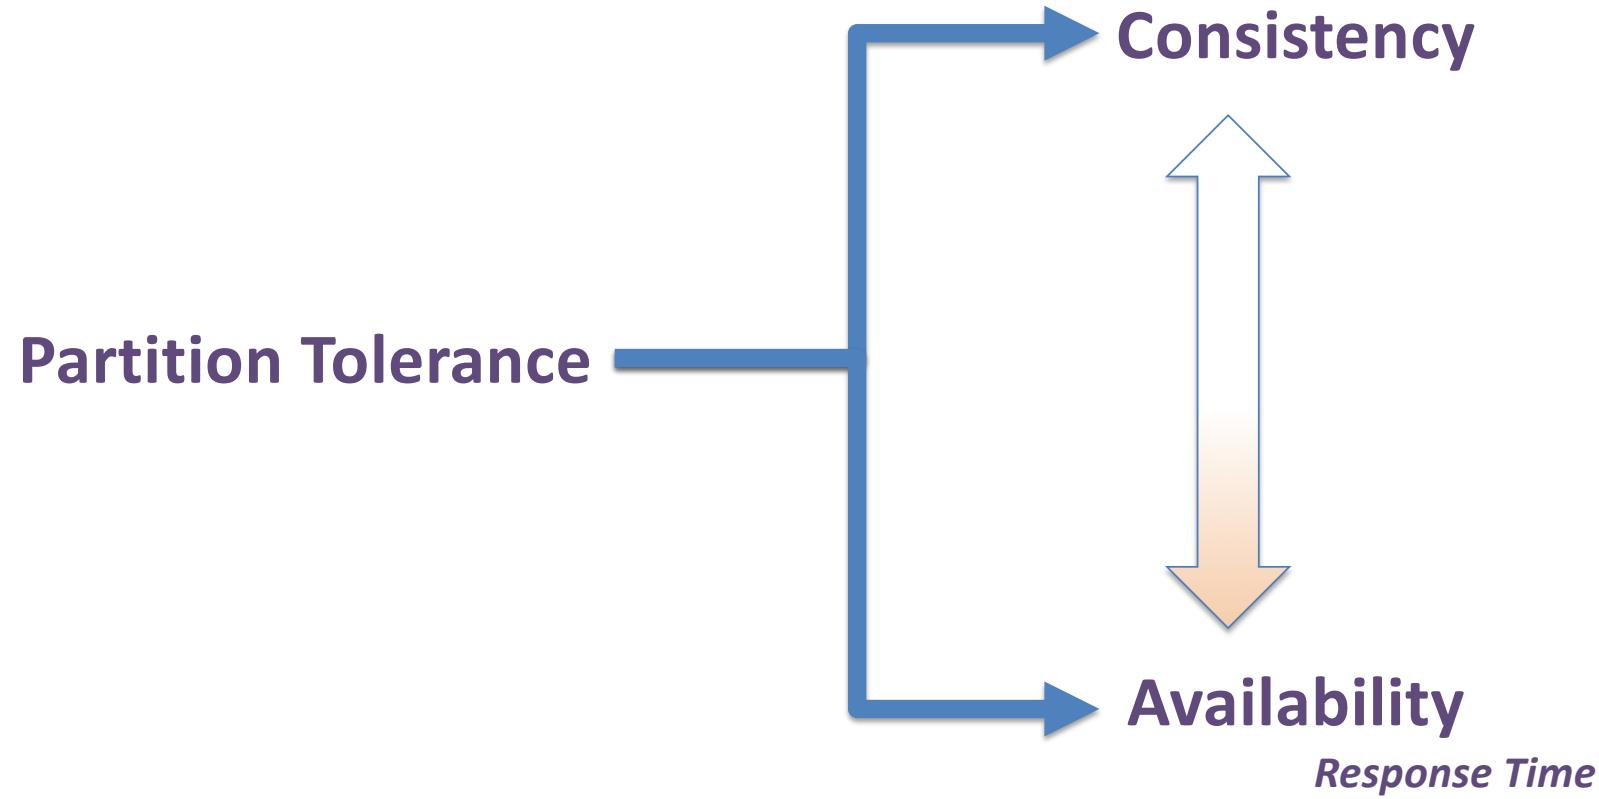
CAP Theorem [Brewer 2000, Lynch 2002]

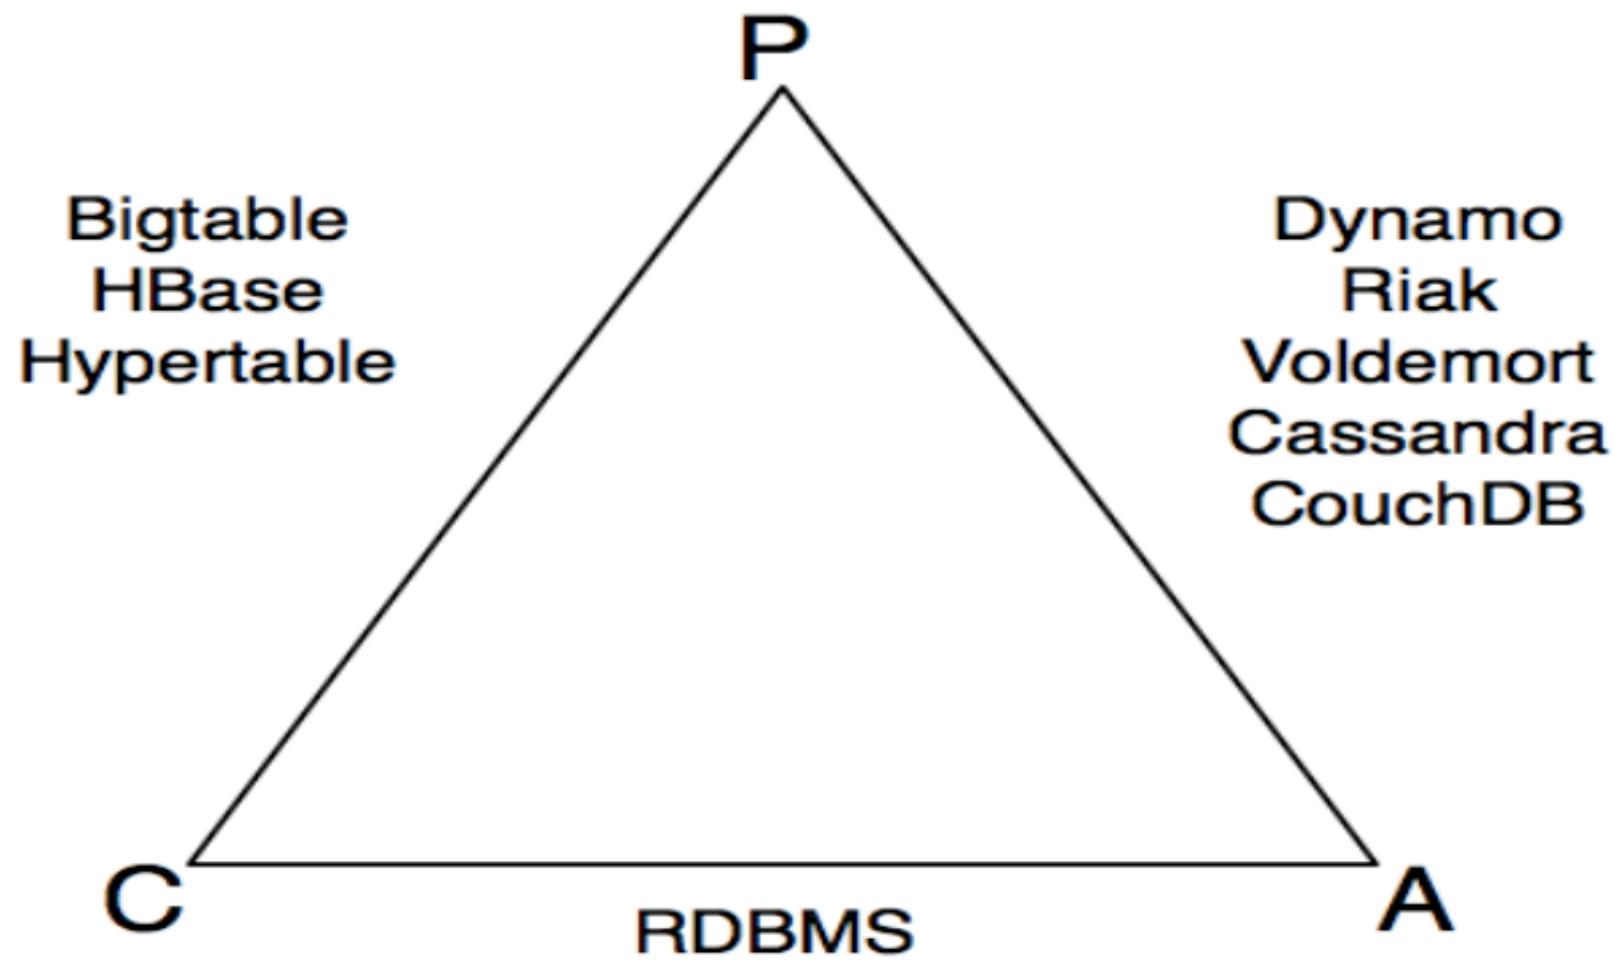
You cannot have
them all.

Availability









src: Shashank Tiwari

Thank you.